

1 Probability

Sum Rule $P(X = x_i) = \sum_{j=1}^J p(X = x_i, Y = y_i)$
 Product rule $P(X, Y) = P(Y|X)P(X)$
 Independence $P(X, Y) = P(X)P(Y)$
 Conditional Independence $P(X, Y|Z) = P(X|Z)P(Y|Z)$
 Bayes' Rule $P(Y|X) = \frac{P(X|Y)P(Y)}{P(X)} = \frac{\sum_{i=1}^k P(X|Y_i)P(Y_i)}{\sum_{i=1}^k P(X|Y_i)}$
 Cond. Ind. $X \perp Y|Z \implies P(X, Y|Z) = P(X|Z)P(Y|Z) = P(X|Z)P(Y|Z)$
 Cond. Ind. $X \perp Y|Z \implies P(X|Y, Z) = P(X|Z)$
 $\mathbb{E}[X] = \int_X t \cdot f_X(t) dt =: \mu_X$
 $\text{Cov}(X, Y) = \mathbb{E}_{x,y}[(X - \mathbb{E}_x[X])(Y - \mathbb{E}_y[Y])]$
 $\text{Cov}(X) := \text{Cov}(X, X) = \text{Var}[X]$
 $X, Y \text{ independent} \implies \text{Cov}(X, Y) = 0$
 $\mathbf{XX}^T \geq 0$ (symmetric positive semidefinite)
 $\text{Var}[X] = \mathbb{E}[X^2] - \mathbb{E}[X]^2$
 $\text{Var}[\mathbf{AX}] = \mathbf{A} \text{Var}[X] \mathbf{A}^T \quad \text{Var}[aX + b] = a^2 \text{Var}[X]$

$$\text{Var} \left[\sum_{i=1}^n a_i X_i \right] = \sum_{i=1}^n a_i^2 \text{Var}[X_i] + 2 \sum_{i < j} a_i a_j \text{Cov}(X_i, X_j)$$

$$\frac{\partial}{\partial t} P(X \leq t) = \frac{\partial}{\partial t} F_X(t) = f_X(t) \quad (\text{derivative of c.d.f. is p.d.f.})$$

$$f_{\alpha Y}(z) = \frac{1}{\alpha} f_Y(\frac{z}{\alpha})$$

T. (Moment Generating Function)

The moment generating function (MGF) $\psi_X(t) = \mathbb{E}[e^{tX}]$ characterizes the distribution of a random variable X .
 $Be(p) = pe^t + (1-p)$
 $N(\mu, \sigma^2) = \exp\left(\mu t + \frac{1}{2}\sigma^2 t^2\right)$
 $Bin(n, p) = (pe^t + (1-p))^n$
 $Gam(\alpha, \beta) = \left(\frac{1}{a - \beta t}\right)$
 for $t < 1/\beta$
 $Pois(\lambda) = e^{\lambda(e^t - 1)}$

T. If X_1, \dots, X_n are indep. rvs with MGFs $M_{X_i}(t) = \mathbb{E}[e^{tX_i}]$, then the MGF of $Y = \sum_{i=1}^n a_i X_i$ is $M_Y(t) = \prod_{i=1}^n M_{X_i}(a_i t)$.
T. Let X, Y be indep., then the p.d.f. of $Z = X + Y$ is the conv. of the p.d.f. of X and Y : $f_Z(z) = \int_{\mathbb{R}} f_X(x) f_Y(z-x) dx$

D. (Normal Distribution)
 $\mathcal{N}(\mathbf{x}; \boldsymbol{\mu}, \boldsymbol{\Sigma}) = \frac{1}{\sqrt{(2\pi)^d \det(\boldsymbol{\Sigma})}} \exp\left(-\frac{1}{2}(\mathbf{x} - \boldsymbol{\mu})^T \boldsymbol{\Sigma}^{-1}(\mathbf{x} - \boldsymbol{\mu})\right)$
 $\hat{\boldsymbol{\mu}} = \frac{1}{n} \sum_{i=1}^n \mathbf{x}_i \quad \hat{\boldsymbol{\Sigma}} = \frac{1}{n} \sum_{i=1}^n (\mathbf{x}_i - \hat{\boldsymbol{\mu}})(\mathbf{x}_i - \hat{\boldsymbol{\mu}})^T$

$$P(\begin{bmatrix} \mathbf{a}_1 \\ \vdots \\ \mathbf{a}_m \end{bmatrix}) = \mathcal{N}(\begin{bmatrix} \mathbf{a}_1 \\ \vdots \\ \mathbf{a}_m \end{bmatrix} \mid \begin{bmatrix} \mathbf{u}_1 \\ \vdots \\ \mathbf{u}_m \end{bmatrix}, \begin{bmatrix} \Sigma_{11} & \cdots & \Sigma_{1m} \\ \vdots & \ddots & \vdots \\ \Sigma_{m1} & \cdots & \Sigma_{mm} \end{bmatrix})$$

$$\mathbf{a}_1, \mathbf{u}_1 \in \mathbb{R}^e, \Sigma_{11} \in \mathbb{R}^{e \times e} \text{ p.s.d.}$$

$$\Sigma_{12} \in \mathbb{R}^{e \times f} \text{ p.s.d.}$$

$$\mathbf{a}_2, \mathbf{u}_2 \in \mathbb{R}^f, \Sigma_{22} \in \mathbb{R}^{f \times f} \text{ p.s.d.}$$

$$\Sigma_{21} \in \mathbb{R}^{f \times e} \text{ p.s.d.}$$

$$P(\mathbf{a}_2 | \mathbf{a}_1 = \mathbf{z}) = \mathcal{N}(\mathbf{a}_2 | \mathbf{u}_2 + \Sigma_{21} \Sigma_{11}^{-1}(\mathbf{z} - \mathbf{u}_1), \Sigma_{22} - \Sigma_{21} \Sigma_{11}^{-1} \Sigma_{12})$$

T. (Chebyshev) Let X be a rv with $\mathbb{E}[X] = \mu$ and variance $\text{Var}[X] = \sigma^2 < \infty$. Then for any $\epsilon > 0$, we have $P(|X - \mu| \geq \epsilon) \leq \frac{\sigma^2}{\epsilon^2}$.

2 Analysis

Log-Trick (Identity): $\nabla_{\theta} [\log(p_{\theta}(\mathbf{x}))] = \frac{\partial}{\partial \mathbf{x}} \log(p_{\theta}(\mathbf{x}))$

T. (Cauchy-Schwarz) $\forall \mathbf{u}, \mathbf{v} \in V: \langle \mathbf{u}, \mathbf{v} \rangle \leq \|\mathbf{u}\| \|\mathbf{v}\|$.

$\forall \mathbf{u}, \mathbf{v} \in V: 0 \leq \langle \mathbf{u}, \mathbf{v} \rangle \leq \|\mathbf{u}\| \|\mathbf{v}\|$.

Special case: $(\sum_i x_i y_i)^2 \leq (\sum_i x_i^2)(\sum_i y_i^2)$.

Special case: $\mathbb{E}[XY]^2 \leq \mathbb{E}[X^2]\mathbb{E}[Y^2]$.

D. (Convex Set) A set $S \subseteq \mathbb{R}^d$ is called *convex* if $\forall \mathbf{x}, \mathbf{x}' \in S, \forall \lambda \in [0, 1]: \lambda \mathbf{x} + (1-\lambda)\mathbf{x}' \in S$.
Com. Any point on the line between two points is within the set. \mathbb{R}^d is convex.

D. (Convex Function) A function $f: S \rightarrow \mathbb{R}$ defined on a convex set $S \subseteq \mathbb{R}^d$ is called *convex* if $\forall \mathbf{x}, \mathbf{x}' \in S, \lambda \in [0, 1]: f(\lambda \mathbf{x} + (1-\lambda)\mathbf{x}') \leq \lambda f(\mathbf{x}) + (1-\lambda)f(\mathbf{x}')$

D. (Convex Function) A function $f: S \subseteq \mathbb{R}^d$ is called *convex* if $\forall \mathbf{x}, \mathbf{x}' \in S, \forall \lambda \in [0, 1]: \lambda \mathbf{x} + (1-\lambda)\mathbf{x}' \in S$.
Com. Any point on the line between two points is within the set. \mathbb{R}^d is convex.

D. (Convex Function) A function $f: S \rightarrow \mathbb{R}$ defined on a convex set $S \subseteq \mathbb{R}^d$ is called *convex* if $\forall \mathbf{x}, \mathbf{x}' \in S, \lambda \in [0, 1]: f(\lambda \mathbf{x} + (1-\lambda)\mathbf{x}') \leq \lambda f(\mathbf{x}) + (1-\lambda)f(\mathbf{x}')$

Com. A function is strictly convex if the line segment between any two points on the graph of the function lies strictly above the graph. This guarantees that there is a unique global minimum.

T. (Properties of Convex Functions)

- $f(y) \geq f(x) + \nabla f(x)^T(y - x)$
- $f'(x) \geq 0$
- Local minima are global minima, strictly convex functions have a unique global minimum
- If f, g are convex then $\alpha f + \beta g$ is convex for $\alpha, \beta \geq 0$
- If f, g are convex then $\max(f, g)$ is convex
- If f is convex and g is convex and non-decreasing then $g \circ f$ is convex

D. (Strongly Convex Function) A function f is μ -strongly convex if it curves up at least as much as a quadratic function with curvature $\mu > 0$. For all x, y :

$$f(y) \geq f(x) + \nabla f(x)^T(y - x) + \frac{\mu}{2} \|y - x\|^2$$

Relation to Optimization:

- Guarantee:** Ensures a unique global minimum exists.
- Convergence:** Gradient Descent on strongly convex (and Lipschitz smooth) functions guarantees a linear convergence rate ($O(c^\ell)$ for some $c < 1$).

Condition Number: The convergence speed depends on the condition number $\kappa = L/\mu$. If κ is large (poor conditioning), convergence slows down.

D. (Condition Number) The condition number $\kappa(A)$ measures the sensitivity of a function's output to small perturbations in the input. For a symmetric positive semi-definite matrix (like the Hessian H of a loss function), it is the ratio of the largest to the smallest eigenvalue:

$$\kappa(H) = \frac{|\lambda_{\max}|}{|\lambda_{\min}|} \geq 1$$

Implications for Optimization:

- Well-conditioned** ($\kappa \approx 1$): The contours of the loss function are nearly spherical. Gradient Descent converges quickly and directly toward the minimum.
- Ill-conditioned** ($\kappa \gg 1$): The contours form narrow, elongated ellipses (steep valleys). Gradient Descent tends to oscillate ("zigzag") across the narrow valley rather than moving down the slope, leading to very slow convergence.

Com. Momentum and Adaptive Learning Rate methods (like Adam) are specifically designed to mitigate the issues caused by high condition numbers.

T. (Taylor-Lagrange Formula)

Given a metric spaces (X, d_X) and (Y, d_Y) , a function $f: X \rightarrow Y$ is called *Lipschitz continuous*, if there exists a real constant $L \in \mathbb{R}_0^+$ (*Lipschitz constant*), such that

$$\forall \mathbf{x}_1, \mathbf{x}_2 \in \mathbb{R}^n: \|f(\mathbf{x}_1) - f(\mathbf{x}_2)\| \leq L \cdot \|\mathbf{x}_1 - \mathbf{x}_2\|.$$

Com. If the objective function is L -smooth a step size of $\eta = 1/L$ guarantees convergence.

D. (Lagrangian Formulation) of arg max _{x, y} $f(x, y)$ s.t. $g(x, y) = c: \mathcal{L}(x, y, \gamma) = f(x, y) - \gamma(g(x, y) - c)$

D. (PL Condition) A differentiable function $f(x)$ with global minimum f^* satisfies the μ -Polyak-Lojasiewicz (PL) condition if there exists a constant $\mu > 0$ such that for all x :

$$\frac{1}{2} \|\nabla f(x)\|^2 \geq \mu(f(x) - f^*)$$

Significance:

- Gradient Dominance:** It implies that the gradient magnitude dominates the suboptimality. If the gradient is small, the function value must be close to the optimal f^* .
- Convergence without Convexity:** The PL condition is weaker than strong convexity (it does not require convexity at all). However, it is sufficient to guarantee a linear convergence rate for Gradient Descent.
- Deep Learning:** Over-parameterized neural networks often satisfy the PL condition in the neighborhood of a minimum, explaining fast convergence despite non-convexity.

Convergence Rate: Gradient Descent with step size $\alpha = 1/L$ (where L is the Lipschitz constant) converges as:

$$f(x_k) - f^* \leq \left(1 - \frac{\mu}{L}\right)^k (f(x_0) - f^*)$$

3 Linear Algebra

Kernels are positive semi-definite matrices.

D. (Positive Semi-Definite Matrix)

A symmetric matrix $\mathbf{A} \in \mathbb{R}^{n \times n}$ is PSD if for all non-zero vectors $\mathbf{x} \in \mathbb{R}^n: \mathbf{x}^T \mathbf{A} \mathbf{x} \geq 0$.

Properties:

- The trace $\text{Tr}(A) \geq 0$ and determinant $\det(A) \geq 0$.
- Cholesky Decomposition exists: $A = LL^T$.

T. (Sylvester Criterion) A $d \times d$ matrix is positive semi-definite if and only if all the upper left $k \times k$ for $k = 1, \dots, d$ have a positive determinant.

Negative definite: $d < 0$ for all odd-sized minors, and $d > 0$ for all even-sized minors

Otherwise: indefinite.

D. (Trace) of $\mathbf{A} \in \mathbb{R}^{n \times n}$ is $\text{Tr}(\mathbf{A}) = \sum_{i=1}^n a_{ii}$.

Properties:

- $\text{Tr}(A) = \sum_i \lambda_i$ (sum of eigenvalues).
- Cyclic property: $\text{Tr}(ABC) = \text{Tr}(BCA) = \text{Tr}(CAB)$.
- Linear: $\text{Tr}(A + B) = \text{Tr}(A) + \text{Tr}(B)$ and $\text{Tr}(cA) = c\text{Tr}(A)$.
- $\text{Tr}(A) = \text{Tr}(A^T)$.

D. (Frobenius Norm ($\|\cdot\|_F$)) The square root of the sum of the absolute squares of its elements.

$$\|A\|_F = \sqrt{\sum_{i,j} |A_{ij}|^2}$$

Properties:

- Relation to Trace: $\|A\|_F = \sqrt{\text{Tr}(A^T A)}$.
- Invariant under orthogonal rotations: $\|Q A\|_F = \|A\|_F$ for orthogonal Q .
- Relation to Singular Values: $\|A\|_F = \sqrt{\sum_i \sigma_i^2}$.

4 Derivatives

4.1 Numerator and Denominator Conventions

Jacobian Layout Convention Com. We use the *numerator-layout*

For a vector-valued function $f: \mathbb{R}^n \rightarrow \mathbb{R}^m$ we define

$$\left(\frac{\partial f}{\partial \mathbf{x}}\right)_{ij} := \frac{\partial f_i}{\partial x_j}, \quad \frac{\partial f}{\partial \mathbf{x}} \in \mathbb{R}^{m \times n}.$$

Hence, gradients of scalar-valued functions are row vectors, and the chain rule takes the form

$$\frac{\partial}{\partial \mathbf{x}} [f(\mathbf{g}(\mathbf{x}))] = \frac{\partial f}{\partial \mathbf{g}} \frac{\partial \mathbf{g}}{\partial \mathbf{x}}.$$

Remark. Sometimes the *denominator-layout* is used, where the Jacobian is defined as $(\mathbf{J}_{ij})_{ij} = \frac{\partial f_j}{\partial x_i} \in \mathbb{R}^{m \times n}$. The two conventions are related by transposition.

4.2 Scalar-by-Vector

Denominator Convention

$\frac{\partial}{\partial \mathbf{x}} [u(\mathbf{x})v(\mathbf{x})] = u(\mathbf{x}) \frac{\partial v(\mathbf{x})}{\partial \mathbf{x}} + v(\mathbf{x}) \frac{\partial u(\mathbf{x})}{\partial \mathbf{x}}$

$\frac{\partial}{\partial \mathbf{x}} [v(\mathbf{x})] = \frac{\partial v(\mathbf{x})}{\partial \mathbf{x}}$

$\frac{\partial}{\partial \mathbf{x}} [\mathbf{r}(\mathbf{x})^T \mathbf{g}(\mathbf{x})] = \frac{\partial \mathbf{r}(\mathbf{x})}{\partial \mathbf{x}} \mathbf{g}(\mathbf{x}) + \mathbf{g}(\mathbf{x}) \frac{\partial \mathbf{r}(\mathbf{x})}{\partial \mathbf{x}}$

$\frac{\partial}{\partial \mathbf{x}} [\mathbf{r}(\mathbf{x})^T \mathbf{A}(\mathbf{x})] = \frac{\partial \mathbf{r}(\mathbf{x})}{\partial \mathbf{x}} \mathbf{A}(\mathbf{x}) + \mathbf{A}(\mathbf{x}) \frac{\$

11.3 Gradient Descent

D. (Gradient Descent (GD)) Iteratively moves parameters θ in the direction of the negative gradient of the loss function $J(\theta)$.

$$\theta_{t+1} = \theta_t - \eta \nabla_\theta J(\theta_t)$$

Com. In Stochastic GD (SGD), the gradient is approximated using a single sample (or mini-batch) to introduce noise and escape local minima. Although, the gradient is unbiased it adds variance, this can help to escape local minima and saddle points.

Com. Gradient Flow can be seen as the numerical integration of the continuous-time ordinary differential equation (ODE) $\dot{x} = -\nabla f(x)$.

Com. For an L -smooth function, GD with a fixed step size $\eta \leq \frac{1}{L}$ is optimal.

D. (Polyak Averaging (Averaged SGD)) Instead of using the final parameter vector θ_T , this method uses the arithmetic mean of the parameters traversed during training.

$$\bar{\theta}_t = \frac{1}{t} \sum_{i=1}^t \theta_i \quad \text{or} \quad \bar{\theta}_t = (1-\beta)\bar{\theta}_{t-1} + \beta\theta_t$$

Benefits:

- It effectively increases the effective batch size and reduces the variance of the estimate.
- Allows the use of larger learning rates (longer steps) while still converging to the optimal solution asymptotically.
- Often achieves the optimal convergence rate of $O(1/t)$ for convex problems.

D. (Learning Rate Condition) (Robbins-Monro Conditions), for Stochastic Gradient Descent (SGD) to guarantee convergence to a local minimum (in non-convex cases) or global minimum (in convex cases), the step size schedule α_t must satisfy two conditions:

1. Explore Forever: The steps must sum to infinity to ensure the algorithm can reach the optimum from any starting point, no matter how far.

$$\sum_{t=1}^{\infty} \alpha_t = \infty$$

2. Decay Fast Enough: The squared steps must sum to a finite value to ensure the variance (noise) of the updates tends to zero, preventing the parameters from oscillating forever around the minimum.

$$\sum_{t=1}^{\infty} \alpha_t^2 < \infty$$

Example: A schedule of $\alpha_t = \frac{1}{t}$ satisfies both, whereas $\alpha_t = \frac{1}{\sqrt{t}}$ satisfies the first but not the second.

D. (Momentum) Accelerates SGD by navigating along the relevant direction and softening oscillations in irrelevant directions. It maintains a velocity vector v (exponential moving average of past gradients).

$$\theta^{t+1} = \theta^t - \eta \nabla J(\theta^t) + \beta(\theta^t - \theta^{t-1})$$

where $\beta \in [0, 1)$ is the momentum term (friction).

D. (Nesterov Accelerated Gradient) A "look ahead" version of GD, also called NAG. It computes the gradient at the *approximate future position* of the parameters rather than the current position.

$$\begin{aligned} \theta^{t+1} &= \theta^t + \beta(\theta^t - \theta^{t-1}) \\ \theta^{t+1} &= \theta^{t-1} - \eta \nabla f(\theta^{t+1}) \end{aligned}$$

D. (Adaptive Learning Rate Methods) These methods adjust the learning rate for each parameter individually, scaling them based on the history of gradient magnitudes.

D. (RMSPProp) Designed to resolve the diminishing learning rates of Adagrad. It uses a decaying average of squared gradients.

$$\begin{aligned} E[g^2]_t &= \beta E[g^2]_{t-1} + (1-\beta)(\nabla J(\theta_t))^2 \\ \theta_{t+1} &= \theta_t - \frac{\alpha}{\sqrt{E[g^2]_t + \epsilon}} \nabla J(\theta_t) \end{aligned}$$

D. (Adam (Adaptive Moment Estimation)) Combines Momentum (first moment m_t) and RMSProp (second moment v_t). It also includes bias correction terms \hat{m}_t, \hat{v}_t to account for initialization at zero.

$$\begin{aligned} m_t &= \beta_1 m_{t-1} + (1-\beta_1) \nabla J(\theta_t) \\ v_t &= \beta_2 v_{t-1} + (1-\beta_2)(\nabla J(\theta_t))^2 \\ \theta_{t+1} &= \theta_t - \frac{\alpha}{\sqrt{v_t + \epsilon}} \hat{m}_t \end{aligned}$$

D. (Muon Optimizer) is an optimizer that takes into account the matrix structure of model parameters. Specifically, while optimizing a loss function $\mathcal{L}(\theta)$ that depends on a weight matrix $W \in \mathbb{R}^{d_{\text{out}} \times d_{\text{in}}}$, at iteration t we follow an update:

$$M_t = \mu M_{t-1} + \nabla_W \mathcal{L}(W_t)$$

P_t is orthogonalize(M_t)

$$W_{t+1} = W_t - \eta \sqrt{\frac{d_{\text{out}}}{d_{\text{in}}}} P_t.$$

Muon is motivated by its ability to increase the scale of "rare directions" that are otherwise ignored during learning

Orthogonalization might require SVD ($U\Sigma V^T$, with unitary $U \in \mathbb{R}^{m \times m}$, rectangular diagonal matrix $\Sigma \in \mathbb{R}^{m \times n}$, and unitary $V \in \mathbb{R}^{n \times n}$)

D. (U-Net (2015)) Designed for Biomedical Image Segmentation (pixel-wise classification).

12.1 Convolutional Layers

D. (Transform) A transform T is a mapping from one function space \mathcal{F} to another function space \mathcal{F}' . So $T: \mathcal{F} \rightarrow \mathcal{F}'$.

D. (Linear Transform) A transform T is linear, if for all functions f, g and scalars α, β , $T(\alpha f + \beta g) = \alpha(Tf) + \beta(Tg)$.

D. (Integral Transform) An *integral transform* is any transform T of the following form

$$(Tf)(u) = \int_{t_1}^{t_2} K(t, u) f(t) dt.$$

Com. The fourier transform is an example of an integral transform.

T. Any integral transform is a linear transform.

D. (Convolution) Given two functions $f, h: \mathbb{R} \rightarrow \mathbb{R}$, their convolution is defined as

$$(f * h)(u) := \int_{-\infty}^{\infty} h(t) f(u-t) dt = \int_{-\infty}^{\infty} h(u-t) f(t) dt$$

Com. Whether the convolution exists depends on the properties of f and h (the integral might diverge). However, a typical use is $f = \text{signal}$, and $h = \text{fast decaying kernel function}$.

T. (Convolution Theorem) Any linear, translation-invariant transformation T can be written as a *convolution* with a suitable h .

T. (Convs are commutative and associative)

T. (Convs are shift-invariant), we define $f_\Delta(t) := f(t + \Delta)$. Then

$$(f_\Delta * h)(u) = (f * h)_\Delta(u)$$

D. (Fourier Transform) The fourier transform of a function f is defined as

$$(\mathcal{F}f)(u) := \int_{-\infty}^{\infty} f(t) e^{-2\pi i u t} dt$$

and its inverse as

$$(\mathcal{F}^{-1}f)(u) := \int_{-\infty}^{\infty} f(t) e^{2\pi i u t} dt$$

Com. Convolutional operators can be efficiently computed with point wise multiplication using the Fourier transform.

$$\mathcal{F}(f * h) = \mathcal{F}f \cdot \mathcal{F}h$$

and then transformed back using the inverse Fourier transform.

$$\mathcal{F}^{-1}(\mathcal{F}f * h) = \mathcal{F}^{-1}(\mathcal{F}f \cdot \mathcal{F}h) = f * h$$

12.2 Discrete Time Convolutions

D. (Discrete Convolution)

For $f, h: \mathbb{Z} \rightarrow \mathbb{R}$, we can define the discrete convolution via

$$(f * h)[u] := \sum_{t=-\infty}^{\infty} f[t] h[u-t] = \sum_{t=-\infty}^{\infty} f[u-t] h[t]$$

Com. Note that the use of rectangular brackets suggests that we're using "arrays" (discrete-time samples).

Com. Typically we use a h with finite support (window size).

D. (Multidimensional Discrete Convolution)

For $f, h: \mathbb{R}^d \rightarrow \mathbb{R}$ we have

$$(f * h)[u_1, \dots, u_d] = \sum_{t_1=-\infty}^{\infty} \dots \sum_{t_d=-\infty}^{\infty} f(t_1, \dots, t_d) h(u_1 - t_1, \dots, u_d - t_d) = \sum_{t_1=-\infty}^{\infty} \dots \sum_{t_d=-\infty}^{\infty} f(u_1 - t_1, \dots, u_d - t_d) h(t_1, \dots, t_d)$$

D. (Discrete Cross-Correlation)

Let $f, h: \mathbb{Z} \rightarrow \mathbb{R}$, then

$$(h * f)[u] := \sum_{t=-\infty}^{\infty} h[t] f[u+t] = \sum_{t=-\infty}^{\infty} h[-t] f[u-t]$$

$$(\bar{h} * f)[u] = (\bar{f} * \bar{h})[u] \quad \text{where } \bar{h}(t) = h(-t).$$

aka "sliding inner product", non-commutative, kernel "flipped over" ($u+t$ instead of $u-t$). If kernel symmetric: cross-correlation = convolution.

12.3 Convolution via Matrices

Represent the input signal, the kernel and the output as vectors. Copy the kernel as columns into the matrix offsetting it by one more very time (gives a band matrix (special case of Toeplitz matrix)). Then the convolution is just a matrix-vector product.

D. (AlexNet (2012)) The breakthrough model that popularized Deep Learning on ImageNet.

Structure: 2 Convolutional layers (with Average Pooling) followed by 3 Fully Connected layers.

Key Features: Introduced the concepts of local receptive fields, shared weights, and spatial subsampling. Used Sigmoid/Tanh activations (pre-ReLU).

D. (Inception Network (2014)) Focused on the effect of network depth using a uniform architecture.

Exploding Gradient: If largest singular value $\sigma_{\max}(U) > 1$.

Philosophy: Replace large filters (e.g., $5 \times 5, 7 \times 7$) with stacks of small 3×3 filters.

Reasoning: Two stacked 3×3 layers have the same receptive field as a 5×5 layer but with fewer parameters and more non-linearities (ReLU between layers).

D. (Valid Padding) Only retain values from windows that are fully-contained within the support of the signal f (see 2D example below) = valid padding

12.5 Backpropagation for Convolutions

D. (Receptive Field I_i^l of x_i^l)

The receptive field I_i^l of node x_i^l is defined as $I_i^l := \{j \mid W_{ij}^l \neq 0\}$ where W^l is the Toeplitz matrix of the convolution at layer l .

Com. Hence, the receptive field of a node x_i^l are just nodes which are connected to it and have a non-zero weight.

Com. One may extend the definition of the receptive field over several layers. The further we go back in layer, the bigger the receptive field becomes due to the nested convolutions. The receptive field may be even the entire image after a few layers. Hence, the convolutions have to be small.

We have $\forall j \neq I_i^l: \frac{\partial \mathcal{R}}{\partial x_j^l} = 0$,

Due to *weight-sharing*, the kernel weight h_j^l is re-used for every unit in the target layer at layer l , so when computing the derivative $\frac{\partial \mathcal{R}}{\partial h_j^l}$ we just build an additive combination of all the derivatives (note that some of them might be zero).

12.9.2 Convolutions in Sequences, NLP & Audio

D. (1D Convolutions) Unlike 2D CNNs for images, sequence modeling (Temporal ConvNets) uses 1D filters that slide over the time axis.

Input: Tensor of shape $(Batch, Length, Channels)$. In NLP, "Channels" are the dimensions of the Word Embeddings.

Function: Captures local temporal dependencies (like n -grams in text) effectively.

Advantage: Highly parallelizable (unlike RNNs which are sequential) and computationally efficient.

D. (Embeddings & NLP) CNNs are often applied on top of pre-trained word embeddings (e.g., Word2Vec, GloVe).

A sentence is represented as a matrix ($Length \times \text{Embedding Dim.}$).

Filters of different widths (e.g., covering 2, 3, or 4 words) act as feature detectors for phrases or specific linguistic patterns (e.g., "very good", "not bad") regardless of their position in the sentence.

D. (LSTM (2014)) Long Short Term Memory models maintain a separate cell state C_t controlled by three gates.

D. (Dilated Convolutions) To handle long sequences without losing resolution (pooling), *dilation* introduces gaps between kernel elements.

Receptive Field: Grows exponentially with the dilation factor d ($1, 2, 4, 8, \dots$), allowing the network to capture long-range dependencies with few layers.

D. (WaveNet (2016)) A deep generative model for raw audio waveforms.

Causal Convolutions: Strict ordering ensures the prediction at time t only depends on samples $x_{<t}$ (cannot see the future).

Dilated Causal Convolutions: Stacks layers with increasing dilation factors. This allows the output neuron to have a receptive field of thousands of timesteps (milliseconds of audio) to generate realistic high-fidelity sound structure.

Skip Connections: Uses residual and parameterized skip connections to speed up convergence and allow training of very deep networks.

D. (GRU): Gated Recurrent Unit models simplify LSTMs by merging Cell/Hidden states and Forget/Input gates.

Ex. input image $m \times n \times c$ (c = number of channels)

K convolution kernels: $p \times q$ (valid padding and stride 1)

output dimensions: $(m-p+1) \times (n-q+1) \times K$

#parameters CNN: $K(pqc + 1)$

#parameters of fully-conn. NN with same number of outputs as CNN:

$mnc((m-p+1)($

16.3 — Spectral Graph Theory

Generalizes convolutions using the Graph Laplacian L (discrete curvature).

D. Laplacian: Defined using the Degree matrix D and Adjacency A .

$$L = D - A \quad \text{or Normalized: } \tilde{L} = I - D^{-\frac{1}{2}}AD^{-\frac{1}{2}}$$

D. Spectral Convolution: Uses the Convolution Theorem via the Graph Fourier Transform (Eigenvectors U of L).

$$x * y = U((U^T x) \odot (U^T y))$$

D. ChebNet: Approximates spectral filters using Chebyshev polynomials T_k to avoid expensive Eigendecomposition ($O(N^3)$). It is strictly K -localized.

$$g_\theta(L) \approx \sum_{k=0}^K \theta_k T_k(\tilde{L})$$

16.4 — Attention GNNs (GAT)

Learns dynamic edge weights α_{ij} instead of static adjacency.

D. Attention Mechanism:

- Score: $e_{ij} = \text{LeakyReLU}(a^T [Wx_i || Wx_j])$
- Normalize: $\alpha_{ij} = \text{softmax}_j(e_{ij}) = \frac{\exp(e_{ij})}{\sum_{k \in \mathcal{N}_i} \exp(e_{ik})}$
- Aggregate: $x'_i = \sigma\left(\sum_{j \in \mathcal{N}_i} \alpha_{ij} Wx_j\right)$

17 Theory

— 17.1 — Neural Tangent Kernel (NTK)

D. Linearized DNN: To analyze non-linear DNNs, we can linearize them around initialization θ_0 using a first-order Taylor expansion:

$$h(\beta)(x) = f(x; \theta_0) + \beta \cdot \nabla_\theta f(x; \theta_0)$$

Here, $\nabla_\theta f(x; \theta_0)$ acts as a fixed, random feature map determined by initialization. The optimization of β becomes a convex problem.

D. NTK Definition: The kernel corresponding to these gradient feature maps is the Neural Tangent Kernel:

$$k(x, \xi) = \langle \nabla_\theta f(x), \nabla_\theta f(\xi) \rangle$$

It encodes the similarity between samples x and ξ based on how much their predictions change when parameters are updated.

D. Infinite Width Limit (NTK Regime): As network width $m \rightarrow \infty$ (under specific "NTK parameter scaling"):

- Deterministic Limit: The initial NTK converges to a deterministic kernel k_n that depends only on the initialization law, not the specific random weights.
- NTK Constancy: The kernel remains constant during training ($\frac{d}{dt} k = 0$). This means the feature map does not evolve ("Lazy Training").
- Linear Dynamics: The training dynamics become identical to Kernel Ridge Regression with kernel k_n . The evolution of outputs follows:

$$\dot{f} = K(\theta)(y - f)$$

17.2 — Bayesian DNNs

D. Bayesian Paradigm: Instead of finding a point estimate θ^* , we compute a posterior distribution $p(\theta|S)$ to capture uncertainty.

$$p(\theta|S) \propto p(S|\theta)p(\theta)$$

Predictions are made by marginalizing over the posterior: $p(y|x) = \int p(y|\theta)p(\theta|S)d\theta$.

D. Langevin Dynamics: Since exact inference is intractable, we use sampling. Langevin dynamics injects noise into Gradient Descent to explore the posterior distribution (sampling from $p(\theta|S)$) rather than collapsing to a minimum.

$$v_{t+1} = (1 - \eta\gamma)v_t - \eta\nabla\tilde{E}(\theta) + 2\gamma\eta\epsilon, \epsilon \sim N(0, I)$$

This mimics a physical system with friction and thermal noise.

— 17.3 — Gaussian Processes (GPs)

D. Infinite Width Equivalence (Neal's Theorem): A single-hidden-layer neural network with infinite width ($m \rightarrow \infty$) and i.i.d. priors on weights converges to a Gaussian Process (GP).

Com. Mechanism: By the Central Limit Theorem, the pre-activations (sums of many independent random variables) become Gaussian.

Com. Result: The network output $f(x)$ is a draw from a GP with mean $\mu(x) = 0$ and a specific kernel $k(x, x')$.

D. Deep GPs: This equivalence extends to deep networks. The kernel is defined recursively:

$$K_l(x, x') = E[\phi(f_{l-1}(x))\phi(f_{l-1}(x'))]$$

where the expectation is taken over the GP of the previous layer f_{l-1} .

— 17.4 — Statistical Learning Theory

D. Generalization Error: The gap between performance on training data (empirical risk) and unseen data (expected risk).

$$\text{Gen}(f) = R[f] - R_{\text{emp}}[f]$$

Classical theory (VC-dimension) predicts overfitting for huge models, but DNNs exhibit Double Descent: test error decreases, rises (at interpolation threshold), and then decreases again as width grows.

D. PAC-Bayesian Bounds: Provides generalization bounds for stochastic classifiers (posterior Q) based on their distance from a prior P (KL-divergence).

$$EQ[R(f)] \leq R_{\text{emp}}(Q) + 2sKL(Q||P) + \ln(2s/\delta)$$

Com. Limitations:

- Oversmoothing:** In deep GCNs, repeated mixing causes all node embeddings to converge to the same value.
- Oversquashing:** Exponentially growing information from distant nodes fails to fit into fixed-size vectors.

18 Chain-Rule and Jacobians for Tensors

D. (k -Dimensional Tensor) $\mathbf{T} \in \mathbb{R}^{d_1 \times d_2 \times \dots \times d_k}$

D. (**Tensor Multiplication**)

$$\mathbf{T} = \underset{\mathbf{P} \in \mathbb{R}^{(a+b)}}{\mathbf{P}} \underset{\mathbf{Q} \in \mathbb{R}^{(b+c)}}{\mathbf{Q}} = \underset{\mathbf{P} \in \mathbb{R}^{r_1 \times \dots \times r_a \times s_1 \times \dots \times s_c}}{\mathbf{P}} \underset{\mathbf{Q} \in \mathbb{R}^{s_1 \times \dots \times s_b \times t_1 \times \dots \times t_c}}{\mathbf{Q}}$$

where each entry of \mathbf{T} is computed as follows:

$$T_{i_1, \dots, i_a, k_1, \dots, k_c} := \sum_{j_1, \dots, j_b} P_{i_1, \dots, i_a, j_1, \dots, j_b} Q_{j_1, \dots, j_b, k_1, \dots, k_c}$$

Note that this is just the sum of the multiplications of two numbers which are in corresponding locations in \mathbf{P} and \mathbf{Q} . Essentially, it's the dot product across the dimensions s_1, \dots, s_b .

Note how this tensor-tensor-multiplication is isomorphic to some matrix-matrix product:

$$T_{i_1, \dots, i_a, k_1, \dots, k_c} := \sum_{j_1, \dots, j_b} P_{i_1, \dots, i_a, j_1, \dots, j_b} Q_{j_1, \dots, j_b, k_1, \dots, k_c}$$

T. (**Tensor Chain Rule**)

$y(D): \mathbb{R}^{d_1 \times d_2} \rightarrow \mathbb{R}^{d_3 \times d_4}, L(y): \mathbb{R}^{d_3 \times d_4} \rightarrow \mathbb{R}$

$$\frac{\partial L}{\partial W} = \frac{\partial L}{\partial W} \times_{d_3, d_4} \frac{\partial y}{\partial W}$$

then we have: $T_{i,j,k,l} = \frac{\partial y_{i,j}}{\partial W_{k,l}}$

19 Generative Models

— 19.1 — Variational Autoencoders (VAEs)

Latent variable models $p(x, z) = p(x|z)p(z)$ where the posterior $p(z|x)$ is intractable. VAEs approximate it using a parametric encoder $q_\phi(z|x)$.

D. (**Evidence Lower Bound (ELBO)**): Since $\ln p(x)$ is intractable, we maximize a lower bound (Jensen's Inequality):

$$\ln p_\theta(x) \geq \mathcal{L}(\theta, \phi; x) = \underbrace{\mathbb{E}_{q_\phi(z|x)} [\ln p_\theta(x|z)]}_{\text{Reconstruction}} - \underbrace{D_{KL}(q_\phi(z|x)||p(z))}_{\text{Regularization}}$$

- Encoder (q_ϕ): Maps input x to latent parameters μ, Σ .
- Decoder (p_θ): Reconstructs x from sampled z .

D. (**Reparameterization Trick**): To backpropagate through the stochastic node $z \sim \mathcal{N}(\mu, \Sigma)$, we move the noise outside:

$$z = \mu + \Sigma^{1/2} \odot \epsilon, \quad \text{where } \epsilon \sim \mathcal{N}(0, I)$$

This makes z a deterministic, differentiable function of ϕ and fixed noise ϵ .

— 19.2 — Factor Analysis

defines a proper probabilistic model of the data ($m \ll n$):

- choose a probability density function p_Z over the latents
- define a conditional probability density function $p_{X|Z}$ over observables
- integrate out the latent variables

$$p_X(x) = \int p_Z(z) p_{X|Z}(x|z) dz$$

- use the Gaussian prior density $z \sim \mathcal{N}(0, I)$, $z \in \mathbb{R}^m$
- add linear observation model for $x \in \mathbb{R}^n$

$$x = \mu + Wz + \eta, \quad \eta \sim \mathcal{N}(0, \Sigma), \quad \Sigma = \text{diag}(\sigma_1^2, \dots, \sigma_n^2)$$

$$x \in \mathbb{R}^n, \quad z \in \mathbb{R}^m, \quad W \in \mathbb{R}^{n \times m}$$

- observational noise η is independent of the latents z
- The induced density is itself normal

$$x \sim \mathcal{N}(\mu, WW^\top + \Sigma), \quad \mu = \frac{1}{s} \sum_{i=1}^s x_i$$

- factors are only identifiable up to orthogonal transformations (rotations, reflections, or permutations) in \mathbb{R}^m , because for any orthogonal matrix $Q \in \mathbb{R}^{m \times m}$ holds:

$$(WQ)(WQ)^\top = WQQ^\top W^\top = WW^\top \quad \text{if } QQ^\top = I$$

- posterior is normal with mean and covariance matrix

$$\mu_{z|x} = W^\top (WW^\top + \Sigma)^{-1}(x - \mu)$$

$$\Sigma_{z|x} = I - W^\top (WW^\top + \Sigma)^{-1}W$$

- Given W , it is thus easy to calc. the posterior over z .

$$\mu_{z|x} \longrightarrow W^\top(x - \mu), \quad \Sigma_{z|x} \longrightarrow 0$$

- MLE with a sample set S : $\mu, W \leftarrow \log p(\mu, W|S)$
- The optimality condition of the i -th column of W is

$$w_i = \rho_i u_i, \quad \rho_i = \max\{0, \sqrt{\lambda_i - \sigma^2}\}.$$

where u_i is the i -th principal eigenvector of the sample covariance matrix and λ_i the corresponding eigenvalue.

- For $\sigma^2 \rightarrow 0$, we recover PCA

— 19.3 — Normalizing Flows

Learns a bijective mapping $f: \mathcal{Z} \rightarrow \mathcal{X}$ from a simple distribution p_Z (e.g., Gaussian) to the complex data distribution p_X . Allows exact likelihood computation.

D. Change of Variables:

$$p_X(x) = p_Z(z) \left| \det \frac{\partial f^{-1}(x)}{\partial z} \right| = p_Z(f^{-1}(x)) |\det J_{f^{-1}}(x)|$$

Or in log-domain (maximizing likelihood):

$$\ln p_X(x) = \ln p_Z(z) - \ln \left| \det \frac{\partial f(z)}{\partial z} \right|$$

D. Coupling Layers (RealNVP): To ensure the Jacobian determinant is computationally cheap, we split variables $x_{1:d}$ and $x_{d+1:D}$:

$$y_{1:d} = x_{1:d}$$

$$y_{d+1:D} = x_{d+1:D} \odot \exp(s(x_{1:d})) + t(x_{1:d})$$

The Jacobian is triangular, so $\det J = \prod \exp(s(x_{1:d}))$.

— 19.4 — Gen. Adversarial Networks (GANs)

A minimax game between a Generator G (creates fakes) and Discriminator D (classifies real vs. fake).

D. (**Minimax Objective**):

$$\min_{\mathcal{F}} \max_{\mathcal{D}} V(D, G) = \mathbb{E}_{x \sim p_{\text{data}}} [\log D(x)] + \mathbb{E}_{z \sim p_z} [\log(1 - D(G(z)))]$$

Optimality:

- Optimal Discriminator:** For a fixed G , $D^*(x) = \frac{p_{\text{data}}(x)}{p_{\text{data}}(x) + p_g(x)}$.
- Global Minimum:** Achieved when $p_g = p_{\text{data}}$. The value is $-\log 2$ (related to Jensen-Shannon Divergence).

Com. Training Issues:

- Vanishing Gradients:** If D is perfect, $\log(1 - D(G(z)))$ saturates. Fix: Train G to maximize $\log D(G(z))$ (Non-Saturating Loss).
- Mode Collapse:** G maps all z to a single plausible x to cheat D .

— 19.5 — Denoising Diffusion Models (DDPM)

Learns to reverse a gradual noising process.

D. (**Forward Process (Fixed)**): Markov chain adding Gaussian noise according to schedule β_t :

$$q(x_t|x_{t-1}) = \mathcal{N}(x_t; \sqrt{1 - \beta_t}x_{t-1}, \beta_t I)$$

Closed form sampling at step t (using $\alpha_t = 1 - \beta_t$ and $\bar{\alpha}_t = \prod \alpha_i$):

$$x_t = \sqrt{\bar{\alpha}_t}x_0 + \sqrt{1 - \bar{\alpha}_t}\epsilon, \quad \epsilon \sim \mathcal{N}(0, I)$$

D. (**Reverse Process (Learned)**): Approximated by a neural network with parameters θ :

$$p_\theta(x_{t-1}|x_t) = \mathcal{N}(x_{t-1}; \mu_\theta(x_t, t), \Sigma_\theta(x_t, t))$$

D. (**Simplified Objective**): Instead of predicting the image mean μ , we predict the noise ϵ added at step t :

$$L_{\text{simple}} = \mathbb{E}_{t, x_0, \epsilon} [\|\epsilon - \epsilon_\theta(\sqrt{\bar{\alpha}_t}x_0 + \sqrt{1 - \bar{\alpha}_t}\epsilon, t)\|^2]$$

19.6 — Autoregressive Models

T. (**AR for Density Estimation**)

Any joint density can be factorized as a product of conditionals

$$p(\mathbf{x}) = \prod_{i=1}^d p(x_i | \mathbf{x}_{<i})$$

Traditional linear AR models in statistics use:

$$p(x_i | \mathbf{x}_{<i}) = \mathcal{N}(x_i; \sum_{k=1}^K \theta_k x_{i-k}, \sigma^2)$$

To obtain richer models, we can model each conditional with a DNN Properties:

- exact likelihood evaluation (fully tractable)
- sampling is sequential over dimensions (can be slow)
- AR structure \rightarrow triangular Jacobian in flows

VAEs & GANs are complicated to learn / not always successful. AR models are simple (look at chainrule $p(x_1, \dots, x_m) = \prod_{i=1}^m p(x_i | x_{1:i-1})$)

20 Ethics

— 20.1 — Robustness

D. (**Adversarial examples**): Classification Perspective Input x , label y , budget ϵ , norm $\|\cdot\|_p$ (usually $p \in \{2, \infty\}$) for each attack type:

- Untargeted: $\|\delta\|_p \leq \epsilon$ and $\hat{y}(x + \delta) \neq y$
- Targeted: $\|\delta\|_p \leq \epsilon$ and $\hat{y}(x + \delta) = t, t \neq y$
- Loss-based: $\max_{\|\delta\|_p \leq \epsilon} \mathcal{L}(f(x + \delta), y)$

Binary Classification: $f(x) = w^\top x + b$, adversarial perturbation pushes x across decision boundary if $y(w^\top x + b) \leq 0$.

D. (**Norms**)

- $\|x\|_p = (\sum_{i=1}^d |x_i|^p)^{1/p}$
- $\|x\|_\infty = \max_{i=1, \dots, d} |x_i|$

D. (**Min L2 adversarial perturbation**): Robustness increases with margin $|w^\top x + b|$ and decreases with $\|w\|_2^2$

$$\delta^* = -\frac{w^\top x + b}{\|w\|_2^2} w, \quad \|\delta^*\|_2 = \frac{|w^\top x + b|}{\|w\|_2}$$

T. (**L_∞ threat model**): If $\|w\|_2 \leq \epsilon$, then $w^\top \delta$ is minimized by choosing $\delta = -\epsilon \text{ sign}(yw)$.

Multiclass: $f_k(x) = w_k^\top x + b_k$, A Perturbation δ is (untargeted) adversarial if it violates at least one inequality:

$$\exists j \neq y: (w_y - w_j)^\top (x + \delta) + (b_y - b_j) \leq 0$$

D. (**Margin to class j ≠ y**): $m_j(x) := (w_y - w_j)^\top x + (b_y - b_j)$

D. (**Nearest Competing Class**): $j^*(x) := \arg\min_{j \neq y} \frac{m_j(x)}{\|w_y - w_j\|_2}$

D. (**Local Linearization**): $J(X) \in \mathbb{R}^{K \times d}$ is the Jacobian with rows $\nabla_x f_k(x)^\top$.

D. (**1st Order Approx.**): $f(x + \delta) \approx f(x) + J(x)^\top \delta$

Fast-grad.-sign-method (FGSM) Attack: max. Loss.

$$\delta = \epsilon \text{ sign}(\nabla_x \mathcal{L}(f(x), y)), x^{\text{adv}} = x + \delta_{\text{FGSM}}$$

Projected Grad. Descent (PGD) Attack: iterative refinement of FGSM with projection back onto threat set.

D. (**(LOFO) via Permutation importance**): Let σ be a random permutation of $\{1, \dots, n\}$

$$\frac{1}{n} \sum_{i=1}^n \ell(y^{(i)}, f(x_{-j}^{(i)}, x_j^{(\sigma(i))})) - \frac{1}{n} \sum_{i=1}^n \ell(y^{(i)}, f(x^{(i)}))$$

There are methods that try to better approximate the idealized marginalization by replacing permutation with conditional resampling.

$$\mathbb{P}(X_j | X_{-j} = x_{-j}^{(i)})$$

Context-dependent Contributions:

D. (**Quantity of Interest (Q(S))**): computed using a subset of variables $X_S \subseteq X$.

D. (**Marginal Contribution of X_j in context of S**):

$$\Delta_j(S) = Q(S \cup \{j\}) - Q(S), \quad S \subseteq \{1, \dots, p\} - \{i\}$$

SHAP and SAGE define the importance of X_j by averaging these marginal contributions over all subsets S , using Shapley weights:

$$\min_f \mathbb{E} \left[\max_{\delta \in S} \ell(Y, f(X + \delta)) \right], \quad S = \{\delta : \|\delta\|_p \leq \epsilon\}$$

T. (**Distribution (P, Q) Shift**): Data P → Deployment Q.

Seek $\sup_{Q \in \mathcal{U}(P)} \mathbb{E}_Q [l(f(Z))]$

Robust statistics studies the stability of statistical procedures under small deviations from an assumed model.

D. (**Huber's contamination model**): for arbitrary contaminating dist. Q .

$$P_\epsilon = (1 - \epsilon)P + \epsilon Q$$

T. (**Distributionally Robust Optim.**) also known as (DRO) seeks to find a model f that is robust to small deviations from an assumed model P .

$$\sup_{Q \in \mathcal{U}(P)} \mathbb{E}_Q [\ell(f(Z))], \quad \mathcal{U}(P) = \text{neighborhood of } P$$

D. (**Wasserstein Balls (around P)**):

$$\mathcal{U}_\epsilon(P) = \{Q : W_p(P, Q) \leq \epsilon\}$$

with the Wasserstein p-distance

D. (**Wasserstein-p Distance**):

$$\sum_{j=1}^p \phi_j = \mathbb{Q}(\{1, \dots, p\}) - \mathbb{Q}(\emptyset)$$

Causal effect: what is the causal effect of x_j on y ?

D. (**Causal ordering**) induces a factorization of the joint distribution with $\text{Pa}(X_j)$ denoting the parents (direct causes) of X_j :

$$P(X_1, \dots, X_p) = \prod_{j=1}^p P(X_j | \text{Pa}(X_j))$$

The joint distribution corresponds to a sequential sampling procedure. This ordering reflects causal, not merely statistical, dependencies.

D. (**Structural equation models (SEM)**): U_j is an exogenous noise variable:

$$X_j = f_j(\text{Pa}(X_j), U_j)$$

Intervention in the generative view replaces the sampling step $X_j \leftarrow x_j$.

Intervention in the SEM view replaces the structural equation $X_j = f_j(\text{Pa}(X_j), U_j)$ by $X_j = x_j$.

Counterfactuals correspond to comparing $Y(X)$ and $Y(\text{do}(X_j = x'_j))$ within the same underlying causal model.

— 20.3 — Fairness

Fair treatment of different groups or segments of the population in machine learning.

D. (**Protected Attribute**): characteristic of an individual for which unequal treatment is considered legally, ethically, or socially unacceptable in decision-making systems e.g. sex or gender, race or ethnicity

D. (**Demographic Parity**): A classifier satisfies demographic parity if $\hat{Y} \perp A$, where \hat{Y} is the predicted label and A is the protected attribute. Equivalently:

$$\mathbb{P}(\hat{Y} = 1 | A = a) = \mathbb{P}(\hat{Y} = 1 | A = a') \quad \forall a, a' \in \mathcal{A}$$

D. (**Equalized Odds**): A Classifier satisfies Equalized Odds if $\hat{Y} \perp \text{A|Y}$, where Y is the true label. Equivalently:

$$\mathbb{P}(\hat{Y} = 1 | A = a, Y =$$