

1 Probability

Sum Rule $P(X = x_i) = \sum_{j=1}^J p(X = x_i, Y = y_i)$
 Product rule $P(X, Y) = P(Y|X)P(X)$
 Independence $P(X, Y) = P(X)P(Y)$
 Bayes' Rule $P(Y|X) = \frac{P(X|Y)P(Y)}{P(X)} = \frac{\sum_{i=1}^K P(X|Y_i)P(Y_i)}{\sum_{i=1}^K P(X|Y_i)}$
 Cond. Ind. $X \perp Y|Z \implies P(\bar{X}, \bar{Y}|Z) = P(X|Z)P(Y|Z)$
 Cond. Ind. $X \perp Y|Z \implies P(X|Y, Z) = P(X|Z)$

$$\mathbb{E}[X] = \int_X t \cdot f_X(t) dt =: \mu_X$$

$$\text{Cov}(X, Y) = \mathbb{E}_{x,y}[(X - \mathbb{E}_x[X])(Y - \mathbb{E}_y[Y])]$$

$$\text{Cov}(X) := \text{Cov}(X, X) = \text{Var}[X]$$

$$X, Y \text{ independent} \implies \text{Cov}(X, Y) = 0$$

$$\mathbf{XX}^T \geq 0 \text{ (symmetric positive semidefinite)}$$

$$\text{Var}[X] = \mathbb{E}[X^2] - \mathbb{E}[X]^2$$

$$\text{Var}[\mathbf{AX}] = \mathbf{A} \text{Var}[\mathbf{X}] \mathbf{A}^T \quad \text{Var}[aX + b] = a^2 \text{Var}[X]$$

$$\text{Var} \left[\sum_{i=1}^n a_i X_i \right] = \sum_{i=1}^n a_i^2 \text{Var}[X_i] + 2 \sum_{i < j} a_i a_j \text{Cov}(X_i, X_j)$$

$$\frac{\partial}{\partial t} P(X \leq t) = \frac{\partial}{\partial t} F_X(t) = f_X(t) \text{ (derivative of c.d.f. is p.d.f.)}$$

$$f_{\alpha Y}(z) = \frac{1}{\alpha} f_Y\left(\frac{z}{\alpha}\right)$$

$$\mathbf{T.} \text{The moment generating function (MGF) } \psi_X(t) = \mathbb{E}[e^{tX}] \text{ characterizes the distr. of a rv}$$

$$Be(p) \quad pe^t + (1-p)$$

$$\mathcal{N}(\mu, \sigma) \quad \exp\left(\mu t + \frac{1}{2}\sigma^2 t^2\right)$$

$$Bin(n, p) \quad (pe^t + (1-p))^n$$

$$Gam(\alpha, \beta) \quad \left(\frac{1}{a - \beta t}\right)^{\alpha}$$

$$\text{for } t < 1/\beta$$

$$Pois(\lambda) \quad e^{\lambda(e^t - 1)}$$

$$\mathbf{T.} \text{If } X_1, \dots, X_n \text{ are indep. rvs with MGFs } M_{X_i}(t) = \mathbb{E}[e^{tX_i}], \text{ then the MGF of } Y = \sum_{i=1}^n a_i X_i \text{ is } M_Y(t) = \prod_{i=1}^n M_{X_i}(a_i t).$$

$$\mathbf{T.} \text{Let } X, Y \text{ be indep., then the p.d.f. of } Z = X + Y \text{ is the conv. of the p.d.f. of } X \text{ and } Y: f_Z(z) = \int_X f_X(x) f_Y(z-x) dx$$

$$\mathcal{N}(\mathbf{x}; \boldsymbol{\mu}, \boldsymbol{\Sigma}) = \frac{1}{\sqrt{(2\pi)^d \det(\boldsymbol{\Sigma})}} \exp\left(-\frac{1}{2}(\mathbf{x} - \boldsymbol{\mu})^\top \boldsymbol{\Sigma}^{-1} (\mathbf{x} - \boldsymbol{\mu})\right)$$

$$\hat{\mu} = \frac{1}{n} \sum_{i=1}^n \mathbf{x}_i \quad \boldsymbol{\Sigma} = \frac{1}{n} \sum_{i=1}^n (\mathbf{x}_i - \hat{\mu})(\mathbf{x}_i - \hat{\mu})^\top$$

$$\mathbf{T.} P(\lfloor \mathbf{a}_1 \rfloor) = N\left(\lfloor \mathbf{a}_1 \rfloor \mid \lfloor \mathbf{u}_1 \rfloor, \lfloor \Sigma_{11} \Sigma_{12} \rfloor\right)$$

$$\mathbf{a}_1, \mathbf{u}_1 \in \mathbb{R}^d, \Sigma_{11} \in \mathbb{R}^{d \times d} \text{ p.s.d.}, \Sigma_{12} \in \mathbb{R}^{d \times d} \text{ p.s.d.}$$

$$\mathbf{a}_2, \mathbf{u}_2 \in \mathbb{R}^d, \Sigma_{22} \in \mathbb{R}^{d \times d} \text{ p.s.d.}, \Sigma_{21} \in \mathbb{R}^{d \times d} \text{ p.s.d.}$$

$$P(\mathbf{a}_2 \mid \mathbf{a}_1 = \mathbf{z}) =$$

$$N(\mathbf{a}_2 \mid \mathbf{u}_2 + \Sigma_{21}\Sigma_{11}^{-1}(\mathbf{z} - \mathbf{u}_1), \Sigma_{22} - \Sigma_{21}\Sigma_{11}^{-1}\Sigma_{12})$$

$$\mathbf{T.} \text{(Chebyshev)} \text{ Let } X \text{ be a rv with } \mathbb{E}[X] = \mu \text{ and variance } \text{Var}[X] = \sigma^2 < \infty. \text{ Then for any } \epsilon > 0, \text{ we have } P(|X - \mu| \geq \epsilon) \leq \frac{\sigma^2}{\epsilon^2}.$$

2 Analysis

$$\text{Log-Trick (Identity): } \nabla_\theta [p_\theta(\mathbf{x})] = p_\theta(\mathbf{x}) \nabla_\theta [\log(p_\theta(\mathbf{x}))]$$

$$\mathbf{T.} \text{(Cauchy-Schwarz)}$$

$$\forall u, v \in V: \langle u, v \rangle \leq \|u\| \|v\|.$$

$$\forall u, v \in V: 0 \leq \langle u, v \rangle \leq \|u\| \|v\|.$$

$$\text{Special case: } (\sum_i x_i y_i)^2 \leq (\sum_i x_i^2)(\sum_i y_i^2).$$

$$\text{Special case: } \mathbb{E}[XY]^2 \leq \mathbb{E}[X^2]\mathbb{E}[Y^2].$$

$$\mathbf{D.} \text{(Convex Set)} \text{ A set } S \subseteq \mathbb{R}^d \text{ is called } \text{convex} \text{ if } \forall \mathbf{x}, \mathbf{x}' \in S, \forall \lambda \in [0, 1]:$$

$$\lambda \mathbf{x} + (1 - \lambda)\mathbf{x}' \leq \lambda f(\mathbf{x}) + (1 - \lambda)f(\mathbf{x}')$$

$$\text{Com. A function is strictly convex if the line segment between any two points on the graph of the function lies strictly above the graph. This guarantees that there is a unique global minimum.}$$

$$\mathbf{T.} \text{(Properties of Convex Functions)}$$

$$\cdot f(y) \geq f(x) + \nabla f(x)^T(y - x)$$

$$\cdot f''(x) \geq 0$$

$$\cdot \text{Local minima are global minima, strictly convex functions have a unique global minimum}$$

$$\cdot \text{If } f, g \text{ are convex then } af + bg \text{ is convex for } \alpha, \beta \geq 0$$

$$\cdot \text{If } f \text{ is convex and } g \text{ is convex then } \max(f, g) \text{ is convex}$$

$$\cdot \text{If } f \text{ is convex and } g \text{ is convex and non-decreasing then } g \circ f \text{ is convex}$$

$$\mathbf{D.} \text{(Strongly Convex Function)} \text{ A function } f \text{ is } \mu\text{-strongly convex if it curves up at least as much as a quadratic function with curvature } \mu > 0. \text{ For all } x, y:$$

$$f(y) \geq f(x) + \nabla f(x)^T(y - x) + \frac{\mu}{2} \|y - x\|^2$$

$$\text{Relation to Optimization:}$$

$$\cdot \text{Guarantee: Ensures a unique global minimum exists.}$$

$$\cdot \text{Convergence: Gradient Descent on strongly convex (and Lipschitz smooth) functions guarantees a linear convergence rate } O(c^\ell) \text{ for some } c < 1.$$

$$\cdot \text{Condition Number: The convergence speed depends on the condition number } \kappa = L/\mu. \text{ If } \kappa \text{ is large (poor conditioning), convergence slows down.}$$

$$\mathbf{D.} \text{(Condition Number)} \text{ The condition number } \kappa(A) \text{ measures the sensitivity of a function's output to small perturbations in the input. For a symmetric positive semi-definite matrix (like the Hessian } H \text{ of a loss function), it is the ratio of the largest to the smallest eigenvalue:}$$

$$\kappa(H) = \frac{|\lambda_{\max}|}{|\lambda_{\min}|} \geq 1$$

Implications for Optimization:

- **Well-conditioned** ($\kappa \approx 1$): The contours of the loss function are nearly spherical. Gradient Descent converges quickly and directly toward the minimum.
- **Ill-conditioned** ($\kappa \gg 1$): The contours form narrow, elongated ellipses (steep valleys). Gradient Descent tends to oscillate ("zigzag") across the narrow valley rather than moving down the slope, leading to very slow convergence.

Com. Momentum and Adaptive Learning Rate methods (like Adam) are specifically designed to mitigate the issues caused by high condition numbers.

T. (Taylor-Lagrange Formula)

$$f(x) = \sum_{k=0}^n \frac{f^{(k)}(x_0)}{k!} (x - x_0)^k + \int_{x_0}^x f^{(n+1)}(t) \frac{(x-t)^{n+1}}{n!} dt$$

$$\mathbf{T.} \text{(Jensen)} \quad f \text{ convex/concave, } \forall i: \lambda_i \geq 0, \sum_{i=1}^n \lambda_i = 1$$

$$f(\sum_{i=1}^n \lambda_i x_i) \leq \sum_{i=1}^n \lambda_i f(x_i)$$

$$\text{Special case: } f(\mathbb{E}[X]) \leq \mathbb{E}[f(X)].$$

D. (L-Lipschitz Continuous Function)

Given two metric spaces (X, d_X) and (Y, d_Y) , a function $f: X \rightarrow Y$ is called *Lipschitz continuous*, if there exists a real constant $L \in \mathbb{R}_+$ (*Lipschitz constant*), such that

$$\forall \mathbf{x}_1, \mathbf{x}_2 \in \mathbb{R}^n: \|f(\mathbf{x}_1) - f(\mathbf{x}_2)\| \leq L \cdot \|\mathbf{x}_1 - \mathbf{x}_2\|.$$

Com. If the objective function is *L-smooth* a step size of $\eta = 1/L$ guarantees convergence.

D. (Lagrangian Formulation)

of arg max $_{x,y} f(x, y)$ s.t. $g(x, y) = c: \mathcal{L}(x, y, \gamma) = f(x, y) - \gamma(g(x, y) - c)$

D. (PL Condition)

A differentiable function $f(x)$ with global minimum f^* satisfies the *μ-Polyak-Lojasiewicz (PL)* condition if there exists a constant $\mu > 0$ such that for all x :

$$\frac{1}{2} \|\nabla f(x)\|^2 \geq \mu(f(x) - f^*)$$

4.5 — Vector-by-Matrix (Generalized Gradient)

$$\frac{\partial}{\partial \mathbf{X}} \mathbf{X} = \mathbf{X}$$

5 Information Theory

D. (Entropy)

Let \mathbf{X} be a random variable distributed according to $p(\mathbf{X})$. Then the entropy of \mathbf{X}

$$H(\mathbf{X}) = \mathbb{E}[-\log(p(\mathbf{X}))] = -\sum_{\mathbf{x} \in \mathcal{X}} p(\mathbf{x}) \log(p(\mathbf{x})) \geq 0.$$

It describes the expected information content $I(\mathbf{X})$ of \mathbf{X} .

D. (Cross-Entropy)

For distributions p and q over a given set \mathcal{X} :

$$H(p, q) = -\sum_{\mathbf{x} \in \mathcal{X}} p(\mathbf{x}) \log(q(\mathbf{x})) = \mathbb{E}_{\mathbf{x} \sim p}[-\log(q(\mathbf{x}))]$$

$$H(X; p, q) = H(X) - H(X)$$

where H uses p .

$$\text{The KL-divergence is } q(\mathbf{x}) = 0 \text{ (absolute continuity). Whenever } p(\mathbf{x}) \text{ is zero the contribution of the corresponding term is interpreted as zero because }$$

$$\lim_{x \rightarrow 0^+} x \log(x) = 0.$$

In ML it is a measure of the amount of information lost, when q (model) is used to approximate p (true).

$$\text{Com. } KL(p, q) = 0 \iff p = q.$$

Com. Note that the KL-divergence is not symmetric!

D. (Trace)

of $\mathbf{A} \in \mathbb{R}^{n \times n}$ is $\text{Tr}(\mathbf{A}) = \sum_{i=1}^n a_{ii}$.

Properties:

$$\cdot \text{Tr}(A) = \sum_i \lambda_i \text{ (sum of eigenvalues).}$$

$$\cdot \text{Cyclic property: } \text{Tr}(ABC) = \text{Tr}(BCA) = \text{Tr}(CAB).$$

$$\cdot \text{Linear: } \text{Tr}(A + B) = \text{Tr}(A) + \text{Tr}(B) \text{ and } \text{Tr}(cA) = c\text{Tr}(A).$$

$$\cdot \text{Tr}(A) = \text{Tr}(A^T).$$

D. (Frobenius Norm ($\|\cdot\|_F$))

The square root of the sum of the absolute squares of its elements.

$$\|\mathbf{A}\|_F = \sqrt{\sum_{i,j} |a_{ij}|^2}$$

Properties:

$$\cdot \text{Relation to Trace: } \|A\|_F = \sqrt{\text{Tr}(A^T A)}.$$

$$\cdot \text{Invariant under orthogonal rotations: } \|Q A\|_F = \|A\|_F \text{ for orthogonal } Q.$$

$$\cdot \text{Relation to Singular Values: } \|A\|_F = \sqrt{\sum_i \sigma_i^2}.$$

11.3 Gradient Descent

D. (Gradient Descent (GD)) Iteratively moves parameters θ in the direction of the negative gradient of the loss function $J(\theta)$.

$$\theta_{t+1} = \theta_t - \eta \nabla_\theta J(\theta_t)$$

Com. In Stochastic GD (SGD), the gradient is approximated using a single sample (or mini-batch) to introduce noise and escape local minima. Although, the gradient is unbiased it adds variance, this can help to escape local minima and saddle points.

Com. Gradient Flow can be seen as the numerical integration of the continuous-time ordinary differential equation (ODE) $\dot{x} = -\nabla f(x)$.

D. (Polyak Averaging (Averaged SGD)) Instead of using the final parameter vector θ_T , this method uses the arithmetic mean of the parameters traversed during training.

$$\bar{\theta}_t = \frac{1}{t} \sum_{i=1}^t \theta_i \quad \text{or} \quad \bar{\theta}_t = (1-\beta)\bar{\theta}_{t-1} + \beta\theta_t$$

Benefits:

- Effectively increases the effective batch size and reduces the variance of the estimate.
- Allows the use of larger learning rates (longer steps) while still converging to the optimal solution asymptotically.
- Often achieves the optimal convergence rate of $O(1/t)$ for convex problems.

D. (Learning Rate (Robbins-Monro Conditions))

For Stochastic Gradient Descent (SGD) to guarantee convergence to a local minimum (in non-convex cases) or global minimum (in convex cases), the step size schedule α_t must satisfy two conditions:

1. Explore Forever: The steps must sum to infinity to ensure the algorithm can reach the optimum from any starting point, no matter how far.

$$\sum_{t=1}^{\infty} \alpha_t = \infty$$

2. Decay Fast Enough: The squared steps must sum to a finite value to ensure the variance (noise) of the updates tends to zero, preventing the parameters from oscillating forever around the minimum.

$$\sum_{t=1}^{\infty} \alpha_t^2 < \infty$$

Example: A schedule of $\alpha_t = \frac{1}{t}$ satisfies both, whereas $\alpha_t = \frac{1}{\sqrt{t}}$ satisfies the first but not the second.

D. (Momentum) Accelerates SGD by navigating along the relevant direction and softening oscillations in irrelevant directions. It maintains a velocity vector v (exponential moving average of past gradients).

$$\theta^{t+1} = \theta^t - \eta \nabla J(\theta^t) + \beta(\theta^t - \theta^{t-1})$$

where $\beta \in [0, 1]$ is the momentum term (friction).

D. (Nesterov Accelerated Gradient (NAG)) A "look-ahead" version of GD. It computes the gradient at the approximate future position of the parameters rather than the current position.

$$\begin{aligned} \theta^{t+1} &= \theta^t + \beta(\theta^t - \theta^{t-1}) \\ \theta^{t+1} &= \theta^{t+1} - \eta \nabla f(\theta^{t+1}) \end{aligned}$$

D. (Adaptive Learning Rate Methods) These methods adjust the learning rate for each parameter individually, scaling them based on the history of gradient magnitudes.

D. (RMSProp) Designed to resolve the diminishing learning rates of Adagrad. It uses a decaying average of squared gradients.

$$\begin{aligned} E[g^2]_t &= \beta E[g^2]_{t-1} + (1-\beta)(\nabla J(\theta_t))^2 \\ \theta_{t+1} &= \theta_t - \frac{\alpha}{\sqrt{E[g^2]_t + \epsilon}} \nabla J(\theta^t) \end{aligned}$$

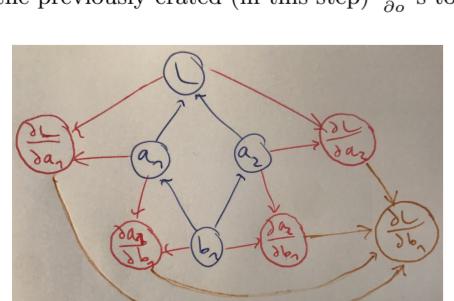
D. (Adam (Adaptive Moment Estimation)) Combines Momentum (first moment m_t) and RMSProp (second moment v_t). It also includes bias correction terms \hat{m}_t, \hat{v}_t to account for initialization at zero.

$$\begin{aligned} m_t &= \beta m_{t-1} + (1-\beta)\nabla J(\theta_t) \\ v_t &= \beta v_{t-1} + (1-\beta)(\nabla J(\theta_t))^2 \\ \theta_{t+1} &= \theta_t - \frac{\alpha}{\sqrt{v_t + \epsilon}} \hat{m}_t \end{aligned}$$

11.4 Backpropagation Graphs

The approach a backpropagation graph for a loss L from a FF network graph is as follows:

1. (RED) Starting at the loss L backward: for each node n and for each of its outputs o which has path to the loss, add the node $\frac{\partial o}{\partial n}$ (at the height of node n). Connect the output o and that node n to that newly created node.
2. (BROWN) Starting at the loss backward: for each node n create a node $\frac{\partial L}{\partial n}$ (if it doesn't exist already) and connect each previously created partial node $(\frac{\partial o}{\partial n})$ to it, and the previously created (in this step) $\frac{\partial L}{\partial o}$'s too.



12.1 Convolutional Layers

D. (Transform) A transform T is a mapping from one function space \mathcal{F} to another function space \mathcal{F}' . So $T: \mathcal{F} \rightarrow \mathcal{F}'$.

D. (Linear Transform) A transform T is linear, if for all functions f, g and scalars α, β , $T(\alpha f + \beta g) = \alpha(Tf) + \beta(Tg)$.

D. (Integral Transform) An integral transform is any transform T of the following form

$$(Tf)(u) = \int_{t_1}^{t_2} K(t, u) f(t) dt.$$

Com. The fourier transform is an example of an integral transform.

T. Any integral transform is a linear transform.

D. (Convolution) Given two functions $f, h: \mathbb{R} \rightarrow \mathbb{R}$, their convolution is defined as

$$(f * h)(u) := \int_{-\infty}^{\infty} h(t)f(u-t) dt = \int_{-\infty}^{\infty} h(u-t)f(t) dt$$

Com. Whether the convolution exists depends on the properties of f and h (the integral might diverge). However, a typical use is $f = \text{signal}$, and $h = \text{fast decaying kernel function}$.

T. (Convolution Theorem) Any linear, translation-invariant transformation T can be written as a convolution with a suitable h .

T. (Convolutions are commutative and associative)

T. (Convolutions are shift-invariant), we define $f_\Delta(t) := f(t + \Delta)$. Then

$$(f_\Delta * h)(u) = (f * h)_\Delta(u)$$

D. (Fourier Transform) The fourier transform of a function f is defined as

$$(\mathcal{F}f)(u) := \int_{-\infty}^{\infty} f(t)e^{-2\pi i u t} dt$$

and its inverse as

$$(\mathcal{F}^{-1}f)(u) := \int_{-\infty}^{\infty} f(t)e^{2\pi i u t} dt$$

Com. Convolutional operators can be efficiently computed with point wise multiplication using the Fourier transform.

$$\mathcal{F}(f * h) = \mathcal{F}^{-1}(\mathcal{F}f \cdot \mathcal{F}h) = f * h$$

12.2 Discrete Time Convolutions

D. (Discrete Convolution)

For $f, h: \mathbb{Z} \rightarrow \mathbb{R}$, we can define the discrete convolution via

$$(f * h)[u] := \sum_{t=-\infty}^{\infty} f[t]h[u-t] = \sum_{t=-\infty}^{\infty} f[u-t]h[t]$$

Com. Note that the use of rectangular brackets suggests that we're using "arrays" (discrete-time samples).

Com. Typically we use a h with finite support (window size).

D. (Multidimensional Discrete Convolution)

For $f, h: \mathbb{R}^d \rightarrow \mathbb{R}$ we have

$$\begin{aligned} (f * h)[u_1, \dots, u_d] &= \sum_{t_1=-\infty}^{\infty} \dots \sum_{t_d=-\infty}^{\infty} f(t_1, \dots, t_d)h(u_1 - t_1, \dots, u_d - t_d) \\ &= \sum_{t_1=-\infty}^{\infty} \dots \sum_{t_d=-\infty}^{\infty} f(u_1 - t_1, \dots, u_d - t_d)h(t_1, \dots, t_d) \end{aligned}$$

D. (Discrete Cross-Correlation)

Let $f, h: \mathbb{Z} \rightarrow \mathbb{R}$, then

$$(h * f)[u] := \sum_{t=-\infty}^{\infty} h[t]f[u+t] = \sum_{t=-\infty}^{\infty} h[-t]f[u-t]$$

$$= (\bar{h} * f)[u] = (\bar{h} * \bar{f})[u] \quad \text{where } \bar{h}(t) = h(-t).$$

aka "sliding inner product", non-commutative, kernel "flipped-over" ($u+t$ instead of $u-t$). If kernel symmetric: cross-correlation = convolution.

12.3 Convolution via Matrices

Represent the input signal, the kernel and the output as vectors. Copy the kernel as columns into the matrix offsetting it by one more very time (gives a band matrix (special case of Toeplitz matrix)). Then the convolution is just a matrix-vector product.

12.4 Border Handling

There are different options to do this

D. (Padding of p) Means we extend the image (or each dimension) by p on both sides ($+2p$) and just fill in a constant there (e.g., zero).

D. (Same Padding) Padding with zeros = same padding ("same" constant, i.e., 0, and we'll get a tensor of the "same" dimensions)

D. (Valid Padding) Only retain values from windows that are fully-contained within the support of the signal f (see 2D example below) = valid padding

12.5 Backpropagation for Convolutions

D. (Receptive Field I_i^l of x_i^l)

The receptive field I_i^l of node x_i^l is defined as $I_i^l := \{j \mid W_{ij}^l \neq 0\}$ where W is the Toeplitz matrix of the convolution at layer l .

Com. Hence, the receptive field of a node x_i^l are just nodes which are connected to it and have a non-zero weight.

Com. One may extend the definition of the receptive field over several layers. The further we go back in layer, the bigger the receptive field becomes due to the nested

convolutions. The receptive field may be even the entire image after a few layers. Hence, the convolutions have to be small.

We have $\forall j \neq I_i^l: \frac{\partial z_i^l}{\partial x_j} = 0$,

D. (Linear Transform) A transform T is linear, if for all functions f, g and scalars α, β , $T(\alpha f + \beta g) = \alpha(Tf) + \beta(Tg)$.

D. (Integral Transform) An integral transform is any transform T of the following form

$$(Tf)(u) = \int_{t_1}^{t_2} K(t, u) f(t) dt.$$

Com. The fourier transform is an example of an integral transform.

T. Any integral transform is a linear transform.

D. (Convolution) Given two functions $f, h: \mathbb{R} \rightarrow \mathbb{R}$, their convolution is defined as

$$(f * h)(u) = \int_{-\infty}^{\infty} h(t)f(u-t) dt = \int_{-\infty}^{\infty} h(u-t)f(t) dt$$

Com. Whether the convolution exists depends on the properties of f and h (the integral might diverge). However, a typical use is $f = \text{signal}$, and $h = \text{fast decaying kernel function}$.

T. (Convolution Theorem) Any linear, translation-invariant transformation T can be written as a convolution with a suitable h .

T. (Convolutions are commutative and associative)

T. (Convolutions are shift-invariant), we define $f_\Delta(t) := f(t + \Delta)$. Then

$$(f_\Delta * h)(u) = (f * h)_\Delta(u)$$

D. (Fourier Transform) The fourier transform of a function f is defined as

$$(\mathcal{F}f)(u) = \int_{-\infty}^{\infty} f(t)e^{-2\pi i u t} dt$$

and its inverse as

$$(\mathcal{F}^{-1}f)(u) = \int_{-\infty}^{\infty} f(t)e^{2\pi i u t} dt$$

Com. Convolutional operators can be efficiently computed with point wise multiplication using the Fourier transform.

$$\mathcal{F}(f * h) = \mathcal{F}^{-1}(\mathcal{F}f \cdot \mathcal{F}h) = f * h$$

D. (Fourier Transform) The fourier transform of a function f is defined as

$$(\mathcal{F}f)(u) = \int_{-\infty}^{\infty} f(t)e^{-2\pi i u t} dt$$

and its inverse as

$$(\mathcal{F}^{-1}f)(u) = \int_{-\infty}^{\infty} f(t)e^{2\pi i u t} dt$$

Com. Convolutional operators can be efficiently computed with point wise multiplication using the Fourier transform.

$$\mathcal{F}(f * h) = \mathcal{F}^{-1}(\mathcal{F}f \cdot \mathcal{F}h) = f * h$$

D. (Fourier Transform) The fourier transform of a function f is defined as

$$(\mathcal{F}f)(u) = \int_{-\infty}^{\infty} f(t)e^{-2\pi i u t} dt$$

and its inverse as

$$(\mathcal{F}^{-1}f)(u) = \int_{-\infty}^{\infty} f(t)e^{2\pi i u t} dt$$

Com. Convolutional operators can be efficiently computed with point wise multiplication using the Fourier transform.

<

- 18.1 — Variational Autoencoders (VAEs)
Latent variable models $p(x, z) = p(x|z)p(z)$ where the posterior $p(z|x)$ is intractable. VAEs approximate it using a parametric encoder $q_\phi(z|x)$.

D. (Evidence Lower Bound (ELBO)): Since $\ln p(x)$ is intractable, we maximize a lower bound (Jensen's Inequality):

$$\ln p_\theta(x) \geq \mathcal{L}(\theta, \phi; x) = \underbrace{\mathbb{E}_{q_\phi(z|x)}[\ln p_\theta(x|z)]}_{\text{Reconstruction}} - \underbrace{D_{KL}(q_\phi(z|x)\|p(z))}_{\text{Regularization}}$$

- Encoder (q_ϕ): Maps input x to latent parameters μ, Σ .
- Decoder (p_θ): Reconstructs x from sampled z .

D. (Reparameterization Trick): To backpropagate through the stochastic node $z \sim \mathcal{N}(\mu, \Sigma)$, we move the noise outside:

$$z = \mu + \Sigma^{1/2} \odot \epsilon, \quad \text{where } \epsilon \sim \mathcal{N}(0, I)$$

This makes z a deterministic, differentiable function of ϕ and fixed noise ϵ .

- 18.2 — Normalizing Flows
Learns a bijective mapping $f : \mathcal{Z} \rightarrow \mathcal{X}$ from a simple distribution p_z (e.g., Gaussian) to the complex data distribution p_x . Allows exact likelihood computation.

D. Change of Variables:

$$p_x(x) = p_z(z) \left| \det \frac{\partial f^{-1}(x)}{\partial x} \right| = p_z(f^{-1}(x)) |\det J_{f^{-1}}(x)|$$

Or in log-domain (maximizing likelihood):

$$\ln p_x(x) = \ln p_z(z) - \ln \left| \det \frac{\partial f(z)}{\partial z} \right|$$

D. Coupling Layers (RealNVP): To ensure the Jacobian determinant is computationally cheap, we split variables $x_{1:d}$ and $x_{d+1:D}$:

$$\begin{aligned} y_{1:d} &= x_{1:d} \\ y_{d+1:D} &= x_{d+1:D} \odot \exp(s(x_{1:d})) + t(x_{1:d}) \end{aligned}$$

The Jacobian is triangular, so $\det J = \prod \exp(s(x_{1:d}))$.

- 18.3 — Generative Adversarial Networks (GANs)

A minimax game between a Generator G (creates fakes) and Discriminator D (classifies real vs. fake).

D. (Minimax Objective):

$$\min_G \max_D V(D, G) = \mathbb{E}_{x \sim p_{data}} [\log D(x)] + \mathbb{E}_{z \sim p_z} [\log(1 - D(G(z)))]$$

Optimality:

- Optimal Discriminator:** For a fixed G , $D^*(x) = \frac{p_{data}(x)}{p_{data}(x) + p_g(x)}$.
- Global Minimum:** Achieved when $p_g = p_{data}$. The value is $-\log 4$ (related to Jensen-Shannon Divergence).

Com. Training Issues:

- Vanishing Gradients:** If D is perfect, $\log(1 - D(G(z)))$ saturates. Fix: Train G to maximize $\log D(G(z))$ (Non-Saturating Loss).
- Mode Collapse:** G maps all z to a single plausible x to cheat D .

- 18.4 — Denoising Diffusion Models (DDPM)

Learns to reverse a gradual noising process.

D. (Forward Process (Fixed)): Markov chain adding Gaussian noise according to schedule β_t :

$$q(x_t | x_{t-1}) = \mathcal{N}(x_t; \sqrt{1 - \beta_t} x_{t-1}, \beta_t I)$$

Closed form sampling at step t (using $\alpha_t = 1 - \beta_t$ and $\bar{\alpha}_t = \prod \alpha_i$):

$$x_t = \sqrt{\bar{\alpha}_t} x_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon, \quad \epsilon \sim \mathcal{N}(0, I)$$

D. (Reverse Process (Learned)): Approximated by a neural network with parameters θ :

$$p_\theta(x_{t-1} | x_t) = \mathcal{N}(x_{t-1}; \mu_\theta(x_t, t), \Sigma_\theta(x_t, t))$$

D. (Simplified Objective): Instead of predicting the image mean μ , we predict the noise ϵ added at step t :

$$L_{simple} = \mathbb{E}_{t, x_0, \epsilon} [\|\epsilon - \epsilon_\theta(\sqrt{\bar{\alpha}_t} x_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon, t)\|^2]$$

$$\epsilon_\theta(\sqrt{\bar{\alpha}_t} x_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon, t)$$