

# 1 Probability

Sum Rule  $P(X = x_i) = \sum_j p(X = x_i, Y = y_i)$   
 Product rule  $P(X, Y) = P(Y|X)P(X)$   
 Independence  $P(X, Y) = P(X)P(Y)$   
 Conditional Independence  $P(X, Y|Z) = P(X|Z)P(Y|Z)$   
 Bayes' Rule  $P(Y|X) = \frac{P(X|Y)P(Y)}{P(X)} = \frac{\sum_i P(X|Y_i)P(Y_i)}{\sum_i P(X|Y_i)}$   
 Cond. Ind.  $X \perp Y|Z \implies P(X, Y|Z) = P(X|Z)P(Y|Z)$   
 Cond. Ind.  $X \perp Y|Z \implies P(X|Y, Z) = P(X|Z)$   
 $E[X] = \int_X t \cdot f_X(t) dt =: \mu_X$   
 $\text{Cov}(X, Y) = E_{x,y}[(X - E_x[X])(Y - E_y[Y])]$   
 $\text{Cov}(X) := \text{Cov}(X, X) = \text{Var}[X]$   
 $X, Y \text{ independent} \implies \text{Cov}(X, Y) = 0$   
 $\mathbf{XX}^T \geq 0$  (symmetric positive semidefinite)  
 $\text{Var}[X] = E[X^2] - E[X]^2$   
 $\text{Var}[\mathbf{AX}] = \mathbf{A} \text{Var}[X] \mathbf{A}^T \quad \text{Var}[aX + b] = a^2 \text{Var}[X]$

**D. (Condition Number)** The condition number  $\kappa(A)$  measures the sensitivity of a function's output to small perturbations in the input. For a symmetric positive semi-definite matrix (like the Hessian  $H$  of a loss function), it is the ratio of the largest to the smallest eigenvalue:

$$\kappa(H) = \frac{|\lambda_{\max}|}{|\lambda_{\min}|} \geq 1$$

## Implications for Optimization:

- Well-conditioned ( $\kappa \approx 1$ ):** The contours of the loss function are nearly spherical. Gradient Descent converges quickly and directly toward the minimum.
- Ill-conditioned ( $\kappa \gg 1$ ):** The contours form narrow, elongated ellipses (steep valleys). Gradient Descent tends to oscillate ("zigzag") across the narrow valley rather than moving down the slope, leading to very slow convergence.
- Com.** Momentum and Adaptive Learning Rate methods (like Adam) are specifically designed to mitigate the issues caused by high condition numbers.

## T. (Taylor-Lagrange Formula)

$$f(x) = \sum_{k=0}^n \frac{f^{(k)}(x_0)}{k!} (x - x_0)^k + \int_{x_0}^x \frac{f^{(n+1)}(t)}{n!} dt$$

$$\mathbf{T. (Jensen)} \quad f \text{ convex/concave}, \forall i: \lambda_i \geq 0, \sum_i^n \lambda_i = 1$$

$$f(\sum_{i=1}^n \lambda_i x_i) \leq / \geq \sum_{i=1}^n \lambda_i f(x_i)$$

Special case:  $f(\mathbb{E}[X]) \leq \mathbb{E}[f(X)]$ .

**D. (L-Lipschitz Continuous Function)** Given two metric spaces  $(X, d_X)$  and  $(Y, d_Y)$ , a function  $f: X \rightarrow Y$  is called *Lipschitz continuous*, if there exists a real constant  $L \in \mathbb{R}_0^+$  (*Lipschitz constant*), such that

$$\forall x_1, x_2 \in \mathbb{R}^n: \|f(x_1) - f(x_2)\| \leq L \cdot \|x_1 - x_2\|.$$

Com. If the objective function is  $L$ -smooth a step size of  $\eta = 1/L$  guarantees convergence.

**D. (Lagrangian Formulation)** of  $\arg \max_{x,y} f(x, y)$  s.t.  $g(x, y) = c: L(x, y, \gamma) = f(x, y) - \gamma(g(x, y) - c)$

**D. (PL Condition)** A differentiable function  $f(x)$  with global minimum  $f^*$  satisfies the  $\mu$ -Polyak-Lojasiewicz (PL) condition if there exists a constant  $\mu > 0$  such that for all  $x$ :

$$\frac{1}{2} \|\nabla f(x)\|^2 \geq \mu(f(x) - f^*)$$

**Significance:**

**Gradient Dominance:** It implies that the gradient magnitude dominates the suboptimality. If the gradient is small, the function value must be close to the optimal  $f^*$ .

**Convergence without Convexity:** The PL condition is weaker than strong convexity (it does not require convexity at all). However, it is sufficient to guarantee a linear convergence rate for Gradient Descent.

**In Deep Learning:** Over-parameterized neural networks often satisfy the PL condition in the neighborhood of a minimum, explaining fast convergence despite non-convexity.

**Convergence Rate:** Gradient Descent with step size  $\alpha = 1/L$  (where  $L$  is the Lipschitz constant) converges as:

$$f(x_k) - f^* \leq \left(1 - \frac{\mu}{L}\right)^k (f(x_0) - f^*)$$

## 3 Linear Algebra

Kernels are positive semi-definite matrices.

**D. (Positive Semi-Definite Matrix)** A symmetric matrix  $\mathbf{A} \in \mathbb{R}^{n \times n}$  is PSD if for all non-zero vectors  $\mathbf{x} \in \mathbb{R}^n: \mathbf{x}^T \mathbf{A} \mathbf{x} \geq 0$  Properties:

- All eigenvalues  $\lambda_i \geq 0$ .
- The trace  $\text{Tr}(\mathbf{A}) \geq 0$  and determinant  $\det(\mathbf{A}) \geq 0$ .
- Cholesky Decomposition exists:  $\mathbf{A} = LL^T$ .

**T. (Sylvester's Criterion)** A  $d \times d$  matrix is positive semi-definite if and only if all the upper left  $k \times k$  for  $k = 1, \dots, d$  have a positive determinant.

Negative definite:  $d < 0$  for all odd-sized minors, and  $d > 0$  for all even-sized minors

Otherwise: indefinite.

**D. (Trace)** of  $\mathbf{A} \in \mathbb{R}^{n \times n}$  is  $\text{Tr}(\mathbf{A}) = \sum_{i=1}^n a_{ii}$ .

Properties:

- $\text{Tr}(\mathbf{A}) = \sum_i \lambda_i$  (sum of eigenvalues).
- Cyclic property:  $\text{Tr}(ABC) = \text{Tr}(BCA) = \text{Tr}(CAB)$ .
- Linear:  $\text{Tr}(A + B) = \text{Tr}(A) + \text{Tr}(B)$  and  $\text{Tr}(c\mathbf{A}) = c\text{Tr}(\mathbf{A})$ .
- $\text{Tr}(\mathbf{A}) = \text{Tr}(\mathbf{A}^T)$ .

**D. (Frobenius Norm)  $(\|\cdot\|_F)$**  The square root of the sum of the absolute squares of its elements.

$$\|A\|_F = \sqrt{\sum_{i,j} |A_{ij}|^2}$$

**Properties:**

- Relation to Trace:  $\|A\|_F = \sqrt{\text{Tr}(A^T A)}$ .
- Invariant under orthogonal rotations:  $\|QA\|_F = \|A\|_F$  for orthogonal  $Q$ .
- Relation to Singular Values:  $\|A\|_F = \sqrt{\sum_i \sigma_i^2}$ .

## 4 Derivatives

### 4.1 Numerator and Denominator Convention

**Jacobian Layout Convention** We use the numerator-layout

For a vector-valued function  $f: \mathbb{R}^n \rightarrow \mathbb{R}^m$  we define

$$\left(\frac{\partial f}{\partial \mathbf{x}}\right)_{ij} := \frac{\partial f_i}{\partial x_j}, \quad \frac{\partial f}{\partial \mathbf{x}} \in \mathbb{R}^{n \times m}.$$

Hence, gradients of scalar-valued functions are row vectors, and the chain rule takes the form

$$\frac{\partial}{\partial \mathbf{x}} [\mathbf{f}(\mathbf{g}(\mathbf{x}))] = \frac{\partial \mathbf{f}}{\partial \mathbf{g}} \frac{\partial \mathbf{g}}{\partial \mathbf{x}}.$$

**Remark.** Sometimes the denominator-layout is used, where the Jacobian is defined as  $(\mathbf{J}_{ij})_{ij} = \frac{\partial f_j}{\partial x_i} \in \mathbb{R}^{m \times n}$ . The two conventions are related by transposition.

**Connection between Sigmoid and Tanh (Equal Representation Strength)**

## 4.2 Scalar-by-Vector

**Denominator Convention**  $\frac{\partial u(\mathbf{x})v(\mathbf{x})}{\partial \mathbf{x}} = u(\mathbf{x})\frac{\partial v(\mathbf{x})}{\partial \mathbf{x}} + v(\mathbf{x})\frac{\partial u(\mathbf{x})}{\partial \mathbf{x}}$

$\frac{\partial}{\partial \mathbf{x}} [u(\mathbf{x})v(\mathbf{x})] = \frac{\partial u(\mathbf{x})}{\partial \mathbf{x}} v(\mathbf{x}) + u(\mathbf{x}) \frac{\partial v(\mathbf{x})}{\partial \mathbf{x}}$

$\frac{\partial}{\partial \mathbf{x}} [\mathbf{f}(\mathbf{x})^\top \mathbf{g}(\mathbf{x})] = \frac{\partial \mathbf{f}(\mathbf{x})}{\partial \mathbf{x}} \mathbf{g}(\mathbf{x}) + \mathbf{f}(\mathbf{x}) \frac{\partial \mathbf{g}(\mathbf{x})}{\partial \mathbf{x}} = \mathbf{J}_f(\mathbf{x})\mathbf{g}(\mathbf{x}) + \mathbf{J}_g(\mathbf{x})\mathbf{f}(\mathbf{x})$

$\frac{\partial}{\partial \mathbf{x}} [\mathbf{f}(\mathbf{x})^\top \mathbf{Ag}(\mathbf{x})] = \frac{\partial \mathbf{f}(\mathbf{x})}{\partial \mathbf{x}} \mathbf{Ag}(\mathbf{x}) + \mathbf{f}(\mathbf{x}) \frac{\partial \mathbf{Ag}(\mathbf{x})}{\partial \mathbf{x}} = (\mathbf{Ab}^\top + \mathbf{b}^\top \mathbf{A})\mathbf{x}$

$\frac{\partial}{\partial \mathbf{x}} [\mathbf{Ax}^\top \mathbf{Ax}] = (\mathbf{A} + \mathbf{A}^\top)\mathbf{x}$

$\frac{\partial}{\partial \mathbf{x}} [(\mathbf{Ax} + \mathbf{b})^\top \mathbf{C}(\mathbf{Dx} + \mathbf{e})] = \mathbf{D}^\top \mathbf{C}^\top (\mathbf{Ax} + \mathbf{b}) + \mathbf{A}^\top \mathbf{C}^\top (\mathbf{Dx} + \mathbf{e})$

$\frac{\partial}{\partial \mathbf{x}} [\|\mathbf{f}(\mathbf{x})\|_2^2] = \frac{\partial}{\partial \mathbf{x}} [\mathbf{f}(\mathbf{x})^\top \mathbf{f}(\mathbf{x})] = 2 \frac{\partial}{\partial \mathbf{x}} [\mathbf{f}(\mathbf{x})] \mathbf{f}(\mathbf{x}) = 2 \tilde{\mathbf{J}} \mathbf{f}(\mathbf{x})$

**D. (Softmax)**  $\text{softmax}(\mathbf{x})_i = f(x_i)$

$\# \text{Regions} = O\left(\left(\frac{n}{d}\right)^{(L-1)d} n^d\right)$

**Com.** Deep networks are exponentially more expressive in terms of complex decision boundaries than shallow networks of the same parameter count.

**T. (Shekhtman (ReLU Basis))** Any continuous piecewise linear function  $g(x)$  on  $[0, 1]$  (a polygonal line) with breakpoints (knots)  $t_1 < \dots < t_k$  can be represented exactly as a weighted sum of ReLU units:

$$g(x) = a + bx + \sum_{i=1}^k c_i \text{ReLU}(x - t_i)$$

**Com.** A single hidden layer of ReLUs acts as a universal basis for 1D splines. This establishes the theoretical link between splines and ReLU networks.

**T. (Baron's Theorem)** For a function  $f$  with a finite Fourier moment  $C_f$  (a measure of smoothness), there exists a neural network  $f_n$  with one hidden layer of  $n$  sigmoidal units such that the integrated squared error is bounded by:

$$\|f - f_n\|_{L_2}^2 \leq \frac{(2C_f)^2}{n}$$

**Com.** The approximation error decays at a rate of  $O(1/\sqrt{n})$ , which is independent of the input dimension  $d$ . This suggests neural networks can overcome the "Curse of Dimensionality" for specific classes of functions.

**T. (Cover)** Capacity of perceptron in  $\mathbb{R}^n$  is  $2n$  random patterns. Max dichotomies of  $S$  (points in gen. pos.) by linear classifier:  $C(s+1, n) = 2 \sum_{i=0}^n \binom{n}{i}$ . Assymptotic Shattering:

$$C(s, n) = \begin{cases} 1 - \mathcal{O}(e^{-n}) & \text{if } s \leq n \\ \frac{1}{2} & \text{if } n < s < 2n \\ \frac{1}{2} \mathcal{O}(e^{-n}) & \text{if } s = 2n \\ \text{otherwise} & \end{cases}$$

**D. (Hopfield Networks)**: Associative Memory. Hebbian Learning: "Neurons that fire together, wire together".

$w_{ij} \propto \sum_i x_i x_j$ . Energy Min.:  $E(\mathbf{x}) = -\frac{1}{2} \mathbf{x}^\top \mathbf{W} \mathbf{x}$ .

**D. (PDP (Rumelhart et al. 1986))**: Parallel Distributed Processing. Introduction of Backpropagation (Generalized Delta Rule). Differentiable activations allow  $\delta$  propagation to hidden layers.

**D. (Residual Networks)** learn an incremental improvement:  $F(x) = x + [\phi(Wx + b)]$ . They utilize "skip connections" through the graph. For every node  $v_i$ , it computes the value  $v_i$  and the derivative  $\delta_i = \frac{\partial v_i}{\partial x}$ .

**DAG Flow:** Derivatives flow in the *same direction* as data (Input  $\rightarrow$  Output). No need to store the full graph in memory.

**Efficiency:** Efficient for functions  $f: \mathbb{R}^n \rightarrow \mathbb{R}^m$  where  $n \ll m$  (few inputs, many outputs).

**Cost:** requires  $n$  passes to compute gradients for  $n$  parameters. In Deep Learning (millions of parameters), this is computationally prohibitive.

**O( $n$ )** passes

**D. (Forward Mode AD)** (Forward Propagation of Tangents) Computes the derivative simultaneously with the function evaluation.

**Mechanism:** Propagates perturbations  $\delta$  forward through the graph. For every node  $v_i$ , it computes the value  $v_i$  and the derivative  $\delta_i = \frac{\partial v_i}{\partial x}$ .

**DAG Flow:** Derivatives flow in the *same direction* as data (Input  $\rightarrow$  Output). No need to store the full graph in memory.

**Assymptotic Shattering:**

$$C(s, n) = \begin{cases} 1 - \mathcal{O}(e^{-n}) & \text{if } s \leq n \\ \frac{1}{2} & \text{if } n < s < 2n \\ \frac{1}{2} \mathcal{O}(e^{-n}) & \text{if } s = 2n \\ \text{otherwise} & \end{cases}$$

**4.3 Vector-by-Vector**

A, C, D, a, b, e not a function of  $\mathbf{x}$ ,

$\mathbf{f} = \mathbf{f}(\mathbf{x}), \mathbf{g} = \mathbf{g}(\mathbf{x}), \mathbf{h} = \mathbf{h}(\mathbf{x}), \mathbf{u} = \mathbf{u}(\mathbf{x}), \mathbf{v} = \mathbf{v}(\mathbf{x})$

$\frac{\partial}{\partial \mathbf{x}} [\mathbf{u}(\mathbf{x})\mathbf{f}(\mathbf{x})] = \mathbf{u}(\mathbf{x})\frac{\partial \mathbf{f}(\mathbf{x})}{\partial \mathbf{x}} + \mathbf{f}(\mathbf{x})\frac{\partial \mathbf{u}(\mathbf{x})}{\partial \mathbf{x}}$

$\frac{\partial}{\partial \mathbf{x}} [\mathbf{a}^\top \mathbf{a}] = \frac{\partial}{\partial \mathbf{x}} \mathbf{a} = \mathbf{a}$

$\frac{\partial}{\partial \mathbf{x}} [\mathbf{x}^\top \mathbf{x}] = 2\mathbf{x}$

$\frac{\partial}{\partial \mathbf{x}} [\mathbf{b}^\top \mathbf{Ax}] = \mathbf{A}^\top \mathbf{b}$

**D. (Learning Rate (Robbins-Monro Conditions))** For Stochastic Gradient Descent (SGD) to guarantee convergence to a local minimum (non-convex cases) or global minimum (in convex cases), the step size schedule  $\alpha_t$  must satisfy two conditions:  
**1. Explore Forever:** The steps must sum to infinity to ensure the algorithm can reach the optimum from any starting point, no matter how far.

$$\sum_{t=1}^{\infty} \alpha_t = \infty$$

**2. Decay Fast Enough:** The squared steps must sum to a finite value to ensure the variance (noise) of the updates tends to zero, preventing the parameters from oscillating forever around the minimum.

$$\sum_{t=1}^{\infty} \alpha_t^2 < \infty$$

**Example:** A schedule of  $\alpha_t = \frac{1}{t}$  satisfies both, whereas  $\alpha_t = \frac{1}{\sqrt{t}}$  satisfies the first but not the second.

**D. (Momentum)** Accelerates SGD by navigating along the relevant direction and softening oscillations in irrelevant directions. It maintains a velocity vector  $v$  (exponential moving average of past gradients).

$$\theta^{t+1} = \theta^t - \eta \nabla J(\theta^t) + \beta(\theta^t - \theta^{t-1})$$

where  $\beta \in [0, 1)$  is the momentum term (friction).

**D. (Nesterov Accelerated Gradient (NAG))** A "look-ahead" version of GD. It computes the gradient at the approximate future position of the parameters rather than the current position.

$$\begin{aligned} \theta^{t+1} &= \theta^t + \beta(\theta^t - \theta^{t-1}) \\ \theta^{t+1} &= \theta^{t+1} - \eta \nabla f(\theta^{t+1}) \end{aligned}$$

**D. (Adaptive Learning Rate Methods)** These methods adjust the learning rate for each parameter individually, scaling them based on the history of gradient magnitudes.

**D. (RMSPROP)** Designed to resolve the diminishing learning rates of Adagrad. It uses a decaying average of squared gradients.

$$\begin{aligned} E[g^2]_t &= \beta E[g^2]_{t-1} + (1-\beta)(\nabla J(\theta_t))^2 \\ \theta_{t+1} &= \theta_t - \frac{\alpha}{\sqrt{E[g^2]_t + \epsilon}} \nabla J(\theta_t) \end{aligned}$$

**D. (Adam (Adaptive Moment Estimation))** Combines Momentum (first moment  $m_t$ ) and RMSProp (second moment  $v_t$ ). It also includes bias correction terms  $\hat{m}_t, \hat{v}_t$  to account for initialization at zero.

$$\begin{aligned} m_t &= \beta m_{t-1} + (1-\beta_1)\nabla J(\theta_t) \\ v_t &= \beta v_{t-1} + (1-\beta_2)(\nabla J(\theta_t))^2 \\ \theta_{t+1} &= \theta_t - \frac{\alpha}{\sqrt{v_t} + \epsilon} \hat{m}_t \end{aligned}$$

## 12 Convolutional Neural Networks

### 12.1 Convolutional Layers

**D. (Transform)** A transform  $T$  is a mapping from one function space  $\mathcal{F}$  to another function space  $\mathcal{F}'$ . So  $T: \mathcal{F} \rightarrow \mathcal{F}'$ .

**D. (Linear Transform)** A transform  $T$  is linear, if for all functions  $f, g$  and scalars  $\alpha, \beta$ ,  $T(\alpha f + \beta g) = \alpha(Tf) + \beta(Tg)$ .

**D. (Integral Transform)** An integral transform is any transform  $T$  of the following form

$$(Tf)(u) = \int_{t_1}^{t_2} K(t, u)f(t) dt.$$

**Com.** The fourier transform is an example of an integral transform.

**T.** Any integral transform is a linear transform.

**D. (Convolution)** Given two functions  $f, h: \mathbb{R} \rightarrow \mathbb{R}$ , their convolution is defined as

$$(f * h)(u) := \int_{-\infty}^{\infty} h(t)f(u-t) dt = \int_{-\infty}^{\infty} h(u-t)f(t) dt$$

**Com.** Whether the convolution exists depends on the properties of  $f$  and  $h$  (the integral might diverge). However, a typical use is  $f$  = signal, and  $h$  = fast decaying kernel function.

**T. (Convolution Theorem)** Any linear, translation-invariant transformation  $T$  can be written as a convolution with a suitable  $h$ .

**T. (Convolutions are commutative and associative)**

**T. (Convolutions are shift-invariant)**, we define  $f_{\Delta}(t) := f(t + \Delta)$ . Then

$$(f_{\Delta} * h)(u) = (f * h)_{\Delta}(u)$$

**D. (Fourier Transform)** The fourier transform of a function  $f$  is defined as

$$(\mathcal{F}f)(u) := \int_{-\infty}^{\infty} f(t)e^{-2\pi i ut} dt$$

and its inverse as

$$(\mathcal{F}^{-1}f)(u) := \int_{-\infty}^{\infty} f(t)e^{2\pi i ut} dt$$

**Com.** Convolutional operators can be efficiently computed with point wise multiplication using the Fourier transform.

$$\mathcal{F}(f * h) = \mathcal{F}f \cdot \mathcal{F}h$$

and then transformed back using the inverse Fourier transform.

$$\mathcal{F}^{-1}(\mathcal{F}(f * h)) = \mathcal{F}^{-1}(\mathcal{F}f \cdot \mathcal{F}h) = f * h$$

### 12.2 Discrete Time Convolutions

#### D. (Discrete Convolution)

For  $f, h: \mathbb{Z} \rightarrow \mathbb{R}$ , we can define the discrete convolution via

$$(f * h)[u] := \sum_{t=-\infty}^{\infty} f[t]h[u-t] = \sum_{t=-\infty}^{\infty} f[u-t]h[t]$$

**Com.** Note that the use of rectangular brackets suggests that we're using "arrays" (discrete-time samples).

**Com.** Typically we use a  $h$  with finite support (window size).

#### D. (Multidimensional Discrete Convolution)

For  $f, h: \mathbb{R}^d \rightarrow \mathbb{R}$  we have

$$\begin{aligned} (f * h)[u_1, \dots, u_d] &= \sum_{t_1=-\infty}^{\infty} \dots \sum_{t_d=-\infty}^{\infty} f(t_1, \dots, t_d)h(u_1 - t_1, \dots, u_d - t_d) \\ &= \sum_{t_1=-\infty}^{\infty} \dots \sum_{t_d=-\infty}^{\infty} f(u_1 - t_1, \dots, u_d - t_d)h[t] \end{aligned}$$

#### D. (Discrete Cross-Correlation)

Let  $f, h: \mathbb{Z} \rightarrow \mathbb{R}$ , then

$$\begin{aligned} (h * f)[u] &:= \sum_{t=-\infty}^{\infty} h[t]f[u+t] = \sum_{t=-\infty}^{\infty} h[-t]f[u-t] \\ &= (\bar{h} * f)[u] = (\bar{h} * f)[u] \quad \text{where } \bar{h}(t) = h(-t). \end{aligned}$$

aka "sliding inner product", non-commutative, kernel "flipped over" ( $u+t$  instead of  $u-t$ ). If kernel symmetric: cross-correlation = convolution.

#### 12.3 Convolution via Matrices

Represent the input signal, the kernel and the output as vectors. Copy the kernel  $n$  columns into the matrix offsetting it by one more every time (gives a band matrix (special case of Toeplitz matrix)). Then the convolution is just a matrix-vector product.

#### 12.4 Border Handling

There are different options to do this

**D. (Padding of  $p$ )** Means we extend the image (or each dimension) by  $p$  on both sides (so  $+2p$ ) and just fill in a constant there (e.g., zero).

**D. (Same Padding)** Padding with zeros = same padding ("same" constant, i.e., 0, and we'll get a tensor of the "same" dimensions)

**D. (Valid Padding)** Only retain values from windows that are fully-contained within the support of the signal  $f$  (see 2D example below) = valid padding

#### 12.5 Backpropagation for Convolutions

##### D. (Receptive Field $I_i^l$ of $x_i^l$ )

The receptive field  $I_i^l$  of node  $x_i^l$  is defined as  $I_i^l := \{j \mid W_{ij}^l \neq 0\}$  where  $W^l$  is the Toeplitz matrix of the convolution at layer  $l$ .

**Com.** Hence, the receptive field of a node  $x_i^l$  are just nodes which are connected to it and have a non-zero weight.

**Com.** One may extend the definition of the receptive field over several layers. The further we go back in layer, the bigger the receptive field becomes due to the nested convolutions. The receptive field may be even the entire image after a few layers. Hence, the convolutions have to be small.

We have  $\forall j \neq I_i^l: \frac{\partial x_i^l}{\partial x_j} = 0$ ,

Due to weight-sharing, the kernel weight  $h_j^l$  is re-used for every unit in the target layer at layer  $l$ , so when computing the derivative  $\frac{\partial \mathcal{R}}{\partial h_j^l}$  we just build an additive combination of all the derivatives (note that some of them might be zero).

$$\frac{\partial \mathcal{R}}{\partial h_j^l} = \sum_{i=1}^{m_l} \frac{\partial \mathcal{R}}{\partial x_i^l} \frac{\partial x_i^l}{\partial h_j^l}$$

**Backpropagations of Convolutions as Convolutions**

$y^{(l)}$  output of  $l$ -th layer  $\mathbf{y}^{(l-1)}$  output of  $(l-1)$ -th layer / input to  $l$ -th layer  $w$  convolution filter  $\frac{\partial \mathcal{R}}{\partial y^{(l)}}$  known

$$\mathbf{y}^{(l+1)} = \mathbf{y}^{(l)} * \mathbf{w}$$

$$\frac{\partial \mathcal{R}}{\partial w_i} = \sum_k \frac{\partial \mathcal{R}}{\partial y^{(l)}_i} \frac{\partial y^{(l)}_i}{\partial w_i} = \sum_k \frac{\partial \mathcal{R}}{\partial y^{(l)}_i} \frac{\partial}{\partial w_i} \left[ \sum_{o=p}^p y^{(l-1)}_o w_o \right] = \sum_k \frac{\partial \mathcal{R}}{\partial y^{(l)}_i} y^{(l-1)}_k$$

$$= \sum_k \frac{\partial \mathcal{R}}{\partial y^{(l)}_i} = \sum_k \frac{\partial \mathcal{R}}{\partial y^{(l)}_i} \text{rot180}(y^{(l-1)})_{k-i}$$

$$= \left( \frac{\partial \mathcal{R}}{\partial y^{(l)}} * \text{rot180}(y^{(l-1)}) \right)_i$$

The derivative  $\frac{\partial \mathcal{R}}{\partial y^{(l)}}$  is analogous.

Note that we just used generalized indices  $i, k, o$  which may be multi-dimensional.

This example omits activation functions and biases, but that could easily be included with the chain-rule.

#### D. (Rotation180)

$$\forall i: \text{rot180}(\mathbf{x})_i = \mathbf{x}_{-i}$$

#### 12.6 Pooling

There are min, max, avg, and softmax pooling. Max pooling is the most frequently used one.

#### D. (Max-Pooling)

- 1D:  $x_i^{\max} = \max\{x_{i+k} \mid 0 \leq k < r\}$
- 2D:  $x_{ij}^{\max} = \max\{x_{i+k, j+l} \mid 0 \leq k, l < r\}$

#### 12.7 Sub-Sampling (aka "Strides")

Often, it is desirable to reduce the size of the feature maps. This can be achieved by skipping some of the input values in the convolution. The stride is the number of steps the kernel takes in each direction.

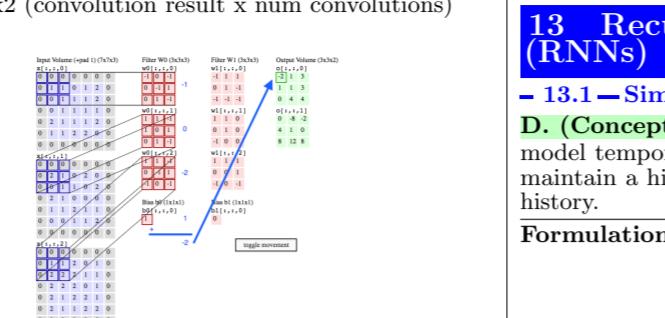
#### 12.8 Channels

Here we have an input signal that is 2D with 3 channels (7x7x3) (image x channels)

- input signal  $m \times n \times c$  ( $c$  = number of channels)
- $K$  convolution kernels:  $p \times q$  (valid padding and stride 1)
- output dimensions:  $(m-p+1) \times (n-q+1) \times K$
- #parameters CNN:  $K(pqc+1)$
- #parameters of fully-conn. NN with same number of outputs as CNN:  $mnc(m-p+1)(n-q+1)+1$

and the Jacobian is triangular, so det  $J = \prod \exp(s_{1:d})$ .

and we want to learn two filters  $W_0$  and  $W_1$ , which each process the 3 channels, and sum the results of the convolutions across each channel leading to a tensor of size  $3 \times 3 \times 2$  (convolution result  $x$  num convolutions)



#parameters of locally-conn. NN with same connections as CNN:  $pqc(m-p+1)(n-q+1)+1K$

**13 Recurrent Neural Networks (RNNs)**

#### 13.1 Simple Recurrent Networks

**D. (Concept)**: Unlike CNNs (fixed filter widths), RNNs model temporal/sequence data of variable length. They maintain a hidden state  $z_t$  acting as a "memory" of the history.

**Formulation:** Given input sequence  $x_1, \dots, x_T$ :

$$z_t = \phi(Uz_{t-1} + Vx_t)$$

$$\hat{y}_t = \psi(Wz_t)$$

where  $U, V, W$  are shared weight matrices and  $\phi, \psi$  are non-linearities.

#### 13.1.4 Sequence Learning

**Generative Modeling:** Decomposes joint probability of sequence:

$$P(x_1, \dots, x_T) = \prod_{t=1}^T P(x_t | x_1, \dots, x_{t-1})$$

**Teacher Forcing:** During training, feed the ground truth  $y_{t-1}^*$  as input to step  $t$  rather than the model's own prediction  $y_{t-1}$ . This speeds up convergence but causes "exposure bias" (train-test discrepancy).

**Seq2Seq (Encoder-Decoder):**

- **Encoder:** Compresses input sequence into a fixed context vector  $z_T$ .
- **Decoder:** Generates output sequence from  $z_T$ .

**Attention:** Allows Decoder to "look back" at specific Encoder states  $z_T$  via learned weights, solving the bottleneck of a fixed-size context vector.

## &lt;h

## 19.3 Generative Adversarial Networks (GANs)

A minimax game between a Generator  $G$  (creates fakes) and Discriminator  $D$  (classifies real vs. fake).

### D. (Minimax Objective):

$$\min_{G} \max_{D} V(D, G) = \mathbb{E}_{x \sim p_{data}} [\log D(x)] + \mathbb{E}_{z \sim p_z} [\log(1 - D(G(z)))]$$

Optimality:

- **Optimal Discriminator:** For a fixed  $G$ ,  $D^*(x) = \frac{p_{data}(x)}{p_{data}(x) + p_g(x)}$ .
- **Global Minimum:** Achieved when  $p_g = p_{data}$ . The value is  $-\log 4$  (related to Jensen-Shannon Divergence).

### Com. Training Issues:

- **Vanishing Gradients:** If  $D$  is perfect,  $\log(1 - D(G(z)))$  saturates. Fix: Train  $G$  to maximize  $\log D(G(z))$  (Non-Saturating Loss).
- **Mode Collapse:**  $G$  maps all  $z$  to a single plausible  $x$  to cheat  $D$ .

## 19.4 Denoising Diffusion Models (DDPM)

Learns to reverse a gradual noising process.

### D. (Forward Process (Fixed)): Markov chain adding Gaussian noise according to schedule $\beta_t$ :

$$q(x_t | x_{t-1}) = \mathcal{N}(x_t; \sqrt{1 - \beta_t} x_{t-1}, \beta_t I)$$

Closed form sampling at step  $t$  (using  $\alpha_t = 1 - \beta_t$  and  $\bar{\alpha}_t = \prod \alpha_i$ ):

$$x_t = \sqrt{\bar{\alpha}_t} x_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon, \quad \epsilon \sim \mathcal{N}(0, I)$$

### D. (Reverse Process (Learned)): Approximated by a neural network with parameters $\theta$ :

$$p_\theta(x_{t-1} | x_t) = \mathcal{N}(x_{t-1}; \mu_\theta(x_t, t), \Sigma_\theta(x_t, t))$$

### D. (Simplified Objective): Instead of predicting the image mean $\mu$ , we predict the noise $\epsilon$ added at step $t$ :

$$L_{simple} = \mathbb{E}_{t, x_0, \epsilon} [\|\epsilon - \epsilon_\theta(\sqrt{\bar{\alpha}_t} x_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon, t)\|^2]$$

## 20 Ethics

### 20.1 Robustness

**D. (Adversarial examples)** (Classification Perspective) Input  $x$ , label  $y$ , budget  $\epsilon$ , norm  $\|\cdot\|_p$  (usually  $p \in [2, \infty)$ ) for each attack type:

- **Untargeted:**  $\|\delta\|_p \leq \epsilon$  and  $\hat{y}(x + \delta) \neq y$
- **Targeted:**  $\|\delta\|_p \leq \epsilon$  and  $\hat{y}(x + \delta) = t, t \neq y$
- **Loss-based:**  $\max_{\|\delta\|_p \leq \epsilon} \mathcal{L}(f(x + \delta), y)$

**Binary Classification:**  $f(x) = w^\top x + b$ , adversarial perturbation pushes  $x$  across decision boundary if  $y(w^\top(x + \delta) + b) \leq 0$ .

### D. (Norms)

- $\|x\|_p = \left( \sum_{i=1}^d |x_i|^p \right)^{1/p}$
- $\|x\|_\infty = \max_{i=1, \dots, d} |x_i|$

**T. (Min L<sub>2</sub> adversarial perturbation):** Robustness increases with margin  $|w^\top x + b|$  and decreases with  $\|w\|_2$

$$\delta^* = \frac{w^\top x + b}{\|w\|_2^2} w, \quad \|\delta^*\|_2 = \frac{|w^\top x + b|}{\|w\|_2}$$

**T. (L<sub>∞</sub> threat model):** If  $\|w\|_2 \leq \epsilon$ , then  $w^\top \delta$  is minimized by choosing  $\delta = -\epsilon \text{ sign}(yw)$ .

**Multiclass:**  $f_k(x) = w_k^\top x + b_k$ . A Perturbation  $\delta$  is (untargeted) adversarial if it violates at least one inequality:

$$\exists j \neq y : (w_y - w_j)^\top (x + \delta) + (b_y - b_j) \leq 0$$

### D. (Margin to class $j \neq y$ ):

$$m_j(x) := (w_y - w_j)^\top x + (b_y - b_j)$$

Nearest Competing Class:  $j^*(x) := \operatorname{argmin}_{j \neq y} \frac{m_j(x)}{\|w_y - w_j\|_2}$

**Neural Networks: Local Linearization** **D. (1st Order Approx.):**  $f(x + \delta) \approx f(x) + J(x)\delta$

$J(X) \in \mathbb{R}^{K \times d}$  is the Jacobian with rows  $\nabla_x f_k(x)^\top$ .

**Fast-grad.-sign-method (FGSM) Attack:** max. Loss.

$$\delta = \epsilon \text{ sign}(\nabla_x \mathcal{L}(f(x), y)), x^{\text{adv}} = x + \delta_{\text{FGSM}}$$

**Projected Grad. Descent (PGD) Attack:** iterative refinement of FGSM with projection back onto threat set.

$$\delta_{t+1} = \operatorname{Proj}_{\|\delta\|_p \leq \epsilon} (\delta_t + \alpha g_t), \quad g_t \in \partial \mathcal{L}(f(x + \delta_t), y).$$

For  $p = \infty$ , a common choice is  $g_t = \text{sign}(\nabla_\delta \mathcal{L}(f(x + \delta_t), y))$ .

**Adversarially robust training:** Instead of  $X$  we evaluate at worst-case loss within neighborhood.

$$\min_{\delta} \mathbb{E}_f [\max_{\delta} \ell(Y, f(X + \delta))], \quad \mathcal{S} = \{\delta : \|\delta\|_p \leq \epsilon\}$$

**T. (Distribution (P, Q) Shift):** Data P → Deployment Q.

Seek  $\sup_{Q \in \mathcal{U}(P)} \mathbb{E}_Q [\ell(f(Z))]$

Robust statistics studies the stability of statistical procedures under small deviations from an assumed model.

**D. (Huber's contamination model):** for arbitrary contaminating dist.  $Q$ .

$$P_\epsilon = (1 - \epsilon)P + \epsilon Q$$

**T. (Distributionally Robust Optimization (DRO)):**

$$\sup_{Q \in \mathcal{U}(P)} \mathbb{E}_Q [\ell(f(Z))], \quad \mathcal{U}(P) = \text{neighborhood of } P$$

**D. (Wasserstein Balls (around P)):**

$$\mathcal{U}_\epsilon(P) = \{Q : W_p(P, Q) \leq \epsilon\}$$

with **D. (Wasserstein-p Distance):**

$$W_p(P, Q) = \left( \inf_{\pi \in \Pi(P, Q)} \int d(z, z')^p d\pi(z, z') \right)^{1/p}$$

For empirical  $P_n = \frac{1}{n} \sum_{i=1}^n \delta_{z_i}$ , inner sup is often deterministic. **T. (Adversarial Transport):**

$$\sup_{Q: W_\infty(P_n, Q) \leq \rho} \left\{ \frac{1}{n} \sum_{i=1}^n g(z'_i) \mid \frac{1}{n} \sum_{i=1}^n d(z'_i, z_i)^p \leq \varepsilon^p \right\}$$

**T. (W<sub>∞</sub> Unification):** Worst case risk over a  $W_\infty$  ball is equivalent to std. adversarial risk with radius  $\rho$ .

$$\sup_{Q: W_\infty(P_n, Q) \leq \rho} \mathbb{E}_Q[g] = \frac{1}{n} \sum_{i=1}^n \sup_{\|x - x_i\|_2 \leq \rho} \ell(f(x), y_i)$$

### 20.2 Interpretability

Interpretability aims to understand the behaviour of a fixed function  $f$  or how the learned function  $f_S$  depends on a subset of variables  $S \subseteq \{1, \dots, p\}$

#### Examples of Local questions:

- *ceteris paribus*: how does the prediction  $f(x)$  change when varying a feature  $x_j$  while keeping all other features fixed?
- $x'_j \mapsto f(x'_j, x_{-j}), \quad x = (x_j, x_{-j})$

**T. (Missing Information):** how does the prediction  $f(x)$  change when a feature  $x_j$  is not observed?

Use **T. (Marginalization)** when the variable  $x_j$  is unobserved, replace the prediction  $f(X)$  with

$$\mathbb{E}[f(X) | X_{-j} = x_{-j}]$$

The contribution of  $x_j$  can be assessed via:  $f(X) - \mathbb{E}[f(X) | X_{-j}]$

**T. (Intervention):** how would the target value change if one could intervene and change the value of a feature  $x_j$ ? **D. (Sensitivity)** is measured by the partial derivative:  $\frac{\partial f(x)}{\partial x_j}$

**T. (Fairness via Attribute-Inv. Latent Rep. Learning)**

We consider the case of equalized odds. Learn latent representation  $Z = f(X)$  s.t. it's predictive of  $Y$  while preventing recovery of the protected attribute  $A$ .

$$X \xrightarrow{f} Z \xrightarrow{h} \hat{Y}, \quad (Z, Y) \xrightarrow{g} \hat{A}$$

This leads to the minimax problem, where  $\ell_{\text{task}}$  and  $\ell_{\text{adv}}$  are classification losses

$$\min_{f, h} \max_g \mathbb{E}[\ell_{\text{task}}(h(Z), Y)] - \lambda \mathbb{E}[\ell_{\text{adv}}(g(Z, Y), A)]$$

**Examples of Global questions:** **Information:** how much info does  $x_j$  carry about  $y$ ?

### D. (Mutual Information):

$$I(X_j; Y) = \mathbb{E} \left[ \log \frac{p(X_j, Y)}{p(X_j)p(Y)} \right]$$

**D. (Conditional Mutual Information):** measures info about  $Y$  uniquely contributed by  $X_j$ , given  $X_{-j}$

$$I(X_j; Y | X_{-j}) = \mathbb{E} \left[ \log \frac{p(X_j, Y | X_{-j})}{p(X_j | X_{-j})p(Y | X_{-j})} \right]$$

**D. (Predictive Utility)** (Leave-one-feature-out (LOFO)): how much risk reduction is obtained by using  $x_j$  (in comb. with other features)?

$$\mathcal{R}(f_{-j}) - \mathcal{R}(f), \quad \mathcal{R}(f_{-j}) = \mathbb{E}[\ell(Y, f_{-j}(X_{-j}))], \quad \mathcal{R}(f) = \dots$$

So for  $f$  and  $f_{-j}$  trained from sufficiently large function classes, the predictive utility is an approximation of the conditional mutual information.

**D. (LOFO via marginalization):** with only one  $f$  trained

$$\mathcal{R}(f_{-j}) = \mathbb{E}[\ell(Y, f(X_{-j}, \tilde{X}_j))], \quad \tilde{X}_j \sim P(X_j | X_{-j})$$

**D. (LOFO via Permutation importance):** Let  $\sigma$  be a random permutation of  $\{1, \dots, n\}$

$$\frac{1}{n} \sum_{i=1}^n \ell(y^{(i)}, f(x_{-j}^{(i)}, x_{\sigma(i)}^{(i)})) - \frac{1}{n} \sum_{i=1}^n \ell(y^{(i)}, f(x^{(i)}))$$

There are methods that try to better approximate the idealized marginalization by replacing permutation with conditional resampling.  $\mathbb{P}(X_j | X_{-j} = x_{-j})$

**Context-dependent Contributions:**

**D. (Quantity of Interest  $\mathcal{Q}(S)$ )** computed using a subset of variables  $X_S \subseteq X$ .

**D. (Marginal Contribution of  $X_j$  in context of  $S$ )**

$$\Delta_j(S) = \mathcal{Q}(S \cup \{j\}) - \mathcal{Q}(S), \quad S \subseteq \{1, \dots, p\} - \{j\}$$

SHAP and SAGE define the importance of  $X_j$  by averaging these marginal contributions over all subsets  $S$ , using Shapley weights:

$$\phi_j = \sum_{S \subseteq \{1, \dots, p\} - \{j\}} \frac{|S|! (p - |S| - 1)!}{p!} (\mathcal{Q}(S \cup \{j\}) - \mathcal{Q}(S))$$

**D. (SHAP (Shapley Additive exPlanations))** For a fixed instance  $x$ ,  $\mathcal{Q}(S)$  is a prediction-level Quantity, attributes predictions:

$$\mathbb{Q}_{\text{SHAP}}(S) = \mathbb{E}[f(X) | X_S = x_S]$$

**D. (SAGE (Shapley Additive Global importance))**

$\mathcal{Q}(S)$  is a performance-level Quantity, attributes risk reduction:

$$\mathbb{Q}_{\text{SAGE}}(S) = -\mathcal{R}(f_S), \quad \mathcal{R}(f_S) = \mathbb{E}[\ell(Y, f_S(X_S))]$$

SHAP and SAGE provide **additive explanations** (each variable is assigned a contribution, and the sum of these contributions recovers the total effect), because Shapley Values  $\{\phi_j\}_{j=1}^p$  satisfy the

**T. (Efficiency (Additivity) Property):**

$$\sum_{j=1}^p \phi_j = \mathcal{Q}(\{1, \dots, p\}) - \mathcal{Q}(\emptyset)$$

**Causal effect:** what is the causal effect of  $x_j$  on  $y$ ?

**D. (Causal ordering)** induces a factorization of the joint distribution with  $\text{Pa}(X_j)$  denoting the parents (direct causes) of  $X_j$ :

$$P(X_1, \dots, X_p) = \prod_{j=1}^p P(X_j | \text{Pa}(X_j))$$

The joint distribution corresponds to a sequential sampling procedure. This ordering reflects causal, not merely statistical, dependencies.

**D. (Structural equation models (SEM))**  $U_j$  is an exogenous noise variable:

$$X_j = f_j(\text{Pa}(X_j), U_j)$$

**Intervention in the generative view** replaces the sampling step  $X_j \leftarrow x_j$ .

**Intervention in the SEM view** replaces the structural equation  $X_j = f_j(\text{Pa}(X_j), U_j)$  by  $X_j = x_j$ .

**Counterfactuals** correspond to comparing  $Y(X)$  and  $Y(\text{do}(X_j = x'_j))$  within the same underlying causal model.